

LATE-TIME COSMOLOGY WITH 21CM INTENSITY MAPPING EXPERIMENTS

PHILIP BULL

Institute of Theoretical Astrophysics, University of Oslo, P.O. Box 1029 Blindern, N-0315 Oslo, Norway

PEDRO G. FERREIRA

Astrophysics, University of Oxford, DWB, Keble Road, Oxford OX1 3RH, UK

PRINA PATEL

Department of Physics, University of Western Cape, Cape Town 7535, South Africa
 SKA SA, 3rd Floor, The Park, Park Road, Pinelands, 7405, South Africa and
 Astrophysics, Cosmology Gravity Centre, and Department of Mathematics and Applied Mathematics, University of Cape Town, Cape Town, 7701, South Africa.

MÁRIO G. SANTOS

Department of Physics, University of Western Cape, Cape Town 7535, South Africa
 SKA SA, 3rd Floor, The Park, Park Road, Pinelands, 7405, South Africa and
 CENTRA, Instituto Superior Técnico, Universidade de Lisboa, Lisboa 1049-001, Portugal
Draft version May 8, 2015

ABSTRACT

We present a framework for forecasting cosmological constraints from future neutral hydrogen intensity mapping experiments at low to intermediate redshifts. In the process, we establish a simple way of comparing such surveys with optical galaxy redshift surveys. We explore a wide range of experimental configurations and assess how well a number of cosmological observables (the expansion rate, growth rate, and angular diameter distance) and parameters (the densities of dark energy and dark matter, spatial curvature, the dark energy equation of state, etc.) will be measured by an extensive roster of upcoming experiments. A number of potential contaminants and systematic effects are also studied in detail. The overall picture is encouraging – if autocorrelation calibration can be controlled to a sufficient level, Phase I of the SKA should be able to constrain the dark energy equation of state about as well as a DETF Stage IV galaxy redshift survey like Euclid, in roughly the same timeframe.

1. INTRODUCTION

As the drive towards ever-greater cosmological precision continues, it becomes necessary to survey progressively larger volumes of the Universe in order to stay ahead of the fundamental limits on measurement accuracy set by cosmic variance. In principle, the best we can ever do is to map out the structure of the whole of the observable Universe, and on large scales at least, this possibility may soon be within reach. The tracer of choice is likely to be neutral hydrogen (HI), which pervades space from the time of recombination up to the present day. HI is thought to be a (biased) tracer of the underlying dark matter distribution, and has a characteristic line emission at around 21cm – well into the radio – that is mostly immune to obscuration by intervening matter. The redshifting of this line additionally gives a measure of cosmic distance, making it possible to reconstruct the three-dimensional matter density field over a wide range of redshifts and scales.

At late times the Universe has reionised, and so the bulk of the neutral hydrogen is thought to reside in comparatively dense gas clouds (damped Lyman-alpha systems) embedded in galaxies, where it is shielded from ionising UV photons. HI is therefore essentially a tracer of the galaxy distribution. Detecting sufficient numbers

of HI-emitting galaxies to do precision cosmology would be a mammoth task, but fortunately this is not necessary; one can instead simply measure the total HI intensity over comparatively large angular scales, without needing to resolve individual galaxies. The result is a map of large-scale fluctuations in 21cm intensity, similar to a CMB map, except now the signal is also a function of redshift. Combined with the high frequency (and thus redshift) resolution of modern radio telescopes, this *intensity mapping* (IM) methodology makes it possible to efficiently survey extremely large volumes (Battye et al. 2004; McQuinn et al. 2006; Chang et al. 2008; Mao et al. 2008; Loeb & Wyithe 2008; Pritchard & Loeb 2008; Wyithe & Loeb 2008; Wyithe et al. 2008; Peterson et al. 2009; Bagla et al. 2010; Seo et al. 2010; Lidz et al. 2011; Ansari et al. 2012; Battye et al. 2013).

As with the CMB, the 21cm signal is contaminated by a host of foreground emission sources, such as our own galaxy and extragalactic point sources, that are orders of magnitude stronger. The hope is that the spectra of the foreground sources are sufficiently smooth that, with a clever cleaning algorithm, it should be possible to suppress them to such a level that the cosmological signal can be extracted in an unbiased way (Oh & Mack 2003; Barkana & Loeb 2005; Santos et al. 2005; Morales et al. 2006; Wang et al. 2006; Gleser et al. 2008; Jelić et al. 2008; Liu et al. 2009; Petrovic & Oh 2011;

Liu & Tegmark 2011; Parsons et al. 2012; Morales et al. 2012; Liu et al. 2014). First attempts at using intensity mapping have been promising, but have highlighted the challenge of calibration and foreground subtraction. The Effelsberg-Bonn survey (Kerp et al. 2011) has produced a data cube covering redshifts out to $z = 0.07$, while the Green Bank Telescope (GBT) has produced the first (tentative) detection of the cosmological signal through IM by cross-correlating with the WiggleZ redshift survey (Chang et al. 2010; Switzer et al. 2013; Masui et al. 2013). As probes to constrain cosmological parameters these measurements are, as yet, ineffective, but they do point the way to a promising future.

The purpose of this paper is two-fold. First of all, we develop a self-consistent forecasting formalism, rooted in the mapping of two-dimensional diffuse emission, but which can easily be compared to (and even interpreted as) 3D redshift surveys of discrete sources. It is approximate – using the “flat-sky” approximation, and slicing up the full dataset into approximately uncorrelated redshift bins – but the formalism is remarkably effective in forecasting constraints for a diverse portfolio of cosmological parameters. We can also use it to discuss the impact of different experimental configurations, as well as the effectiveness of foreground subtraction on our results.

With this formalism in hand, we then use it to explore the observational campaigns that are planned up to, and including, Phase I of the Square Kilometre Array (SKA), which is due to see first light around 2020. One of the key results of this work is the prediction that cosmological constraints from forthcoming 21cm IM surveys will be able to compete with, and perhaps even surpass, those from traditional probes of large scale structure within the decade, even when future high-precision experiments such as Euclid and LSST are taken into account. This finding relies on being able to use large dish arrays like the SKA in an autocorrelation mode, rather than as a (more standard) interferometer – a requirement that brings with it a number of calibration and data analysis challenges that have yet to be solved.

The paper is structured as follows. We first present a mathematical model for the IM signal, and use it to construct the Fisher Matrix for a general set of cosmological parameters. In doing so, we discuss the approximations that we are making in modelling the cosmological signal, experimental set-up, and foreground subtraction. We then discuss the structure of our formalism, comparing it to what one obtains when forecasting for redshift surveys. The concept of “effective volume” becomes a useful way of discussing the strengths and weaknesses of IM. We then make a first pass at forecasting for cosmological parameters, focusing on the detectability of the Baryon Acoustic Oscillation (BAO) feature and the information that can be gleaned from it, including the Hubble rate and angular diameter distance, and the growth rate from redshift space distortions. Next, we turn our attention to forecasting for the canonical set of cosmological parameters, including various fractional energy densities and the equation of state of dark energy. The potential to constrain theories of modified gravity is also examined, and the importance of bias and the neutral hydrogen fraction is assessed. We then discuss the importance of survey design and foreground subtraction, and establish a set of goals that subtraction methods will need to satisfy if IM

experiments are to be successful. We finally conclude with a series of desiderata for the future of cosmology with HI intensity mapping.

Throughout this paper we use the Planck best-fit Λ CDM model (Planck Collaboration 2013b) for our fiducial cosmology,

$$h = 0.67, \quad \Omega_\Lambda = 0.684, \quad \Omega_K = 0, \quad \Omega_b = 0.049, \\ w = -1, \quad n_s = 0.962, \quad \sigma_8 = 0.834, \quad N_{\text{eff}} = 3.046,$$

and all distances and scales are expressed in physical, rather than h^{-1} , units.

2. FISHER FORECAST FORMALISM

We base our forecasting formalism on the *Fisher matrix* technique, which assumes that all parameters of interest can be approximated as being Gaussian-distributed, and that observations are unbiased. While it has its drawbacks (Hawken et al. 2012; Wolz et al. 2012), Fisher forecasting is an effective way of getting an idea of how constraining a given experimental setup is likely to be without requiring detailed experiment-specific simulations.

We begin by defining our data model. The observed brightness temperature is $T_{\text{obs}} = \bar{T}(1 + \delta T)$, where the total fluctuation in an individual “voxel” (volume element) is given by

$$\delta T(\boldsymbol{\theta}_p, \nu_p) = \delta T^S(\boldsymbol{\theta}_p, \nu_p) + \delta T^N(\boldsymbol{\theta}_p, \nu_p) + \delta T^F(\boldsymbol{\theta}_p, \nu_p)$$

with p labelling the voxel given by a 2D angular direction, $\boldsymbol{\theta}_p$, and frequency, ν_p . The total fluctuation consists of the cosmological signal (S), instrumental and atmospheric noise (N), and residual astrophysical foregrounds (F). Detailed models for each component will be set out in subsequent sections, but for now we need only note that they are all stochastic, and will be modelled as Gaussian-distributed, with mean zero, in each voxel.

It is usual to expand the different components of the data vector in terms of spherical harmonics. This is the preferred strategy for accurate forecasting on the largest scales – for example, when testing for scale-dependent bias due to non-Gaussianity (Camera et al. 2013) or effects due to GR/modified gravity (Hall et al. 2013). In this paper, however, we will work in the flat-sky limit, and describe the signal in terms of the comoving 3D power spectrum, $P(\mathbf{k})$. This mimics what is used when forecasting for galaxy redshift surveys, and will be useful when we define an equivalence between redshift surveys and IM experiments.

In this limit, the mapping between the observed voxel and its comoving-space location is¹

$$\mathbf{r}_\perp \approx r(z_i)(\boldsymbol{\theta}_p - \boldsymbol{\theta}_i) \\ r_\parallel \approx \frac{c(1 + z_i)^2}{H(z_i)}(\tilde{\nu}_p - \tilde{\nu}_i) \equiv r_\nu(z_i)(\tilde{\nu}_p - \tilde{\nu}_i),$$

where we have centred the survey on $(\boldsymbol{\theta}_i, \nu_i)$, corresponding to a redshift bin centred at z_i , and have defined the dimensionless frequency $\tilde{\nu} \equiv \nu/\nu_{21} = (1 + z)^{-1}$. We will predominantly work in observational coordinates, with

¹We use the definition of transverse comoving distance from Hogg (1999), which reduces to $r(z) = \int_0^z \frac{cdz}{H(z)}$ for $\Omega_K = 0$, and take $r_\parallel \approx \frac{dr}{dz} dz$.

the Fourier transform convention

$$\delta T(\mathbf{q}, y) = \int \delta T(\boldsymbol{\theta}, \nu) e^{i(\boldsymbol{\theta} \cdot \mathbf{q} + y \cdot \tilde{\nu})} d^2\theta d\tilde{\nu}.$$

The (dimensionless) observation-space Fourier variables are related to the comoving variables by $\mathbf{q} = \mathbf{k}_\perp r$ and $y = k_\parallel r_\nu$.

To construct the Fisher matrix we now need to define the covariance for each of the components. For a component X , this is defined as

$$\langle \delta T^{X*}(\mathbf{q}, y) \delta T^{X'}(\mathbf{q}', y') \rangle = (2\pi)^3 C^X(\mathbf{q}, y) \delta^2(\mathbf{q} - \mathbf{q}') \delta(y - y') \delta_{XX'}. \quad (1)$$

We will make a number of approximations here. First of all, we assume that the signal, noise and foregrounds are uncorrelated with one other, and that the resulting covariance matrices are diagonal (i.e. they are statistically homogeneous and isotropic). The former is not true in practice, as the process by which foregrounds are removed from the data will introduce correlations, as discussed in Section 6.3. The latter is also not strictly true, as the cosmological signal is only diagonal in the flat-sky limit, and the main foreground that we need to correct for – the Galaxy – is anisotropic and will also introduce off-diagonal terms. Nevertheless, given the conservative modelling choices we make in coming sections, we believe our results will be close enough to the real situation.

A further approximation is that evolution can be neglected within each redshift bin, so that evolving cosmological functions are fixed to their values at the central redshift of the bin. This is a good approximation for sufficiently narrow bins, as most of the relevant functions (e.g. $H(z)$, $r(z)$) vary slowly with z . We have verified that our results are robust to the choice of bin width (which is chosen as $\Delta\nu = 60$ MHz for all experiments).

2.1. Signal model

Radio telescopes measure flux density – the integral of the source intensity, I_ν , over the solid angle of the telescope beam. We derive an expression for the HI line intensity in Appendix A. In the Rayleigh-Jeans limit, this can be converted into an effective HI brightness temperature, $T_b = c^2 I_\nu / 2k_B \nu^2$, that can be split into a homogeneous part and a fluctuating part, $T_b = \bar{T}_b(1 + \delta_{\text{HI}})$, where (from Appendix A)

$$\bar{T}_b = \frac{3}{32\pi} \frac{hc^3 A_{10}}{k_B m_p \nu_{21}^2} \frac{(1+z)^2}{H(z)} \Omega_{\text{HI}}(z) \rho_{c,0}. \quad (2)$$

The fluctuations are the quantity of interest here, and so we identify the cosmological signal as

$$\delta T^S(\boldsymbol{\theta}_p, \nu_p) = \bar{T}_b(z) \delta_{\text{HI}}(\mathbf{r}_p, z).$$

At late times, most of the neutral hydrogen content of the Universe is expected to be localised to dense gas clouds within galaxies, where it is shielded from ionising photons. We therefore expect HI to be a biased tracer of the dark matter distribution, just as galaxies are. This allows us to write the HI density contrast as $\delta_{\text{HI}} = b_{\text{HI}} \star \delta_M$ (where δ_M is the total matter density perturbation, and \star denotes convolution, accounting for the possibility of scale- and time-dependent biasing).

Because the HI intensity is measured as a function of frequency (and thus redshift) rather than comoving distance, we must also account for redshift space distortions (RSDs) caused by the peculiar velocities of the clouds and the galaxies in which they reside. Following Kaiser (1987), we write the (Fourier-transformed) redshift-space HI contrast as

$$\delta_{\text{HI}}(\mathbf{k}) = (b_{\text{HI}} + f\mu^2) \exp(-k^2 \mu^2 \sigma_{\text{NL}}^2 / 2) \delta_M(\mathbf{k}), \quad (3)$$

where $\mu \equiv k_\parallel / k$ and the flat-sky approximation has been used again. We have assumed that the HI velocities are unbiased. The linear growth factor, f , is a key observable, telling us much about the growth of structure on linear scales; we will study it in detail in Section 4.3. The exponential term accounts for the ‘‘Fingers of God’’ effect due to uncorrelated velocities on small scales², which washes out structure in the radial direction past a cutoff scale parametrised by the non-linear dispersion, σ_{NL} .

Substituting (3) into (1) and making use of the definition of the isotropic matter power spectrum,

$$\langle \delta_M^*(\mathbf{k}) \delta_M(\mathbf{k}') \rangle \equiv (2\pi)^2 P(k) \delta^3(\mathbf{k} - \mathbf{k}'),$$

we can write the signal covariance as

$$C^S(\mathbf{q}, y) = T_b^2(z_i) \frac{P_{\text{tot}}(z_i, \frac{\mathbf{q}}{r}, \frac{y}{r_\nu})}{r^2 r_\nu}, \quad (4)$$

where the factor of $r^2 r_\nu$ is from the conversion into observational Fourier coordinates, (\mathbf{q}, y) , and we have defined

$$P_{\text{tot}}(z_i, \mathbf{k}_\perp, k_\parallel) = F_{\text{RSD}}(\mathbf{k}_\perp, k_\parallel) D^2(z) P(k, z=0) \\ F_{\text{RSD}}(\mathbf{k}_\perp, k_\parallel) = (b_{\text{HI}} + f\mu^2)^2 \exp(-k^2 \mu^2 \sigma_{\text{NL}}^2).$$

The redshift dependence of the matter power spectrum has been factored out into the linear growth factor, D , which is normalised to $D(z=0) = 1$. This is related to the growth rate by $f = d \log D / d \log a$. Strictly, the growth factor should be scale-dependent on small scales, but for simplicity we neglect this possibility here as we will mostly be concerned with larger scales.

It is straightforward to calculate fiducial values for the cosmological functions in (4). We use CAMB (Lewis et al. 2000) to calculate $P(k)$ at $z=0$ for our chosen fiducial cosmological parameters, and use the simple parametrisation $f(z) = \Omega_M^\gamma(z)$ for the linear growth rate (Peebles 1980; Linder 2005), where $\Omega_M(z) = \Omega_M(1+z)^{3H_0^2/H^2(z)}$, and $\gamma \approx 0.55$ for Λ CDM. For the other functions in (4), however, there is considerably more uncertainty in the choice of fiducial model.

One key uncertainty is the behaviour of the HI bias, b_{HI} . The bias depends on the size of host dark matter haloes; if a halo is too small, gas clouds would be unable to gain sufficient density to shield themselves and keep the hydrogen neutral. The halo dependence can be modelled using the halo mass function with an appropriate lower mass cutoff (or lower rotation velocity); see (Bagla et al. 2010) for example. There are a few candidate models for the evolution of the bias as a function of redshift that fit the current constraints from observations (Switzer et al. 2013) or are calibrated against

²Alternatively, we could have used a Lorentzian instead of an exponential, or a slightly more complex exponential term that models non-linear smoothing of the BAO as well (Samushia et al. 2012).

simulations (Wilman et al. 2008), but there is considerable disagreement between them. In Section 6.1 and Appendix B, we discuss the impact of the uncertain bias evolution. Unless stated otherwise, we will use a linear bias model for the rest of the paper, and – rather conservatively – marginalise over the value of b_{HI} separately in each redshift bin.

Another major uncertainty is in the HI density fraction, $\Omega_{\text{HI}} = \rho_{\text{HI}}/\rho_{\text{c},0}$. This enters the signal covariance through $\bar{T}_b(z)$, since $\bar{T}_b \propto \Omega_{\text{HI}}$. The current best constraints on the HI fraction come from Switzer et al. (2013), who find

$$\Omega_{\text{HI}} b_{\text{HI}} = 4.3 \pm 1.1 \times 10^{-4}$$

at the 68% confidence level at $z = 0.8$. This constitutes a relatively large uncertainty in the overall amplitude of the HI signal and, correspondingly, the signal-to-noise that can be achieved by a given experiment. We investigate the impact of this uncertainty in Section 6.1, but for the rest of the paper we will adopt a fiducial value of $\Omega_{\text{HI},0} = 4.86 \times 10^{-4}$.

The non-linear dispersion scale, σ_{NL} , is yet another source of uncertainty. Recent values from the literature vary between $\sigma_{\text{NL}} \approx 4 - 10$ Mpc (e.g. Li et al. (2007); Reid & White (2011); Reid et al. (2012); Contreras et al. (2013)); for our fiducial model, we choose a middling value of $\sigma_{\text{NL}} = 7$ Mpc (Li et al. 2007), which corresponds to a non-linear scale of $k_{\text{NL}} \sim 0.14 \text{ Mpc}^{-1}$ (or a velocity dispersion of ~ 500 km/s). We check the sensitivity of our results to this choice in Section 6.2. IM experiments can independently constrain σ_{NL} , so we leave it free, marginalising over it as a nuisance parameter in all forecasts.

2.2. Noise model and effective beams

The noise covariance models the instrumental and sky noise for a given experiment, but we will also use it to include the effects of instrumental beams. For radio telescopes, assuming uncorrelated Gaussian noise, the noise covariance has the standard form

$$C^N(\mathbf{q}, y) = \frac{T_{\text{sys}}^2}{t_{\text{tot}} \Delta\nu} U_{\text{bin}} \mathcal{I} B_{\perp}^{-2} B_{\parallel}^{-1}, \quad (5)$$

where T_{sys} is the system temperature, t_{tot} is the total integration time, $U_{\text{bin}} = S_{\text{area}} \Delta\tilde{\nu}$ is the volume of an individual redshift bin, and S_{area} and $\Delta\tilde{\nu}$ are the survey area and (dimensionless) bandwidth within the redshift bin respectively. The factors of \mathcal{I} and B describe the number (or number density) of receivers and their corresponding frequency and angular responses.

The system temperature has two main components: the instrument temperature, T_{inst} , which depends on the hardware design, and $T_{\text{sky}} \approx 60 \text{ K} \times (\nu/300 \text{ MHz})^{-2.5}$, which accounts for atmospheric and background radio emission, to give a total $T_{\text{sys}} = T_{\text{inst}} + T_{\text{sky}}$. Values for T_{inst} are quoted in the design specifications for a given experiment, and are typically a few tens of Kelvin.

The survey area and total integration time are not intrinsic to the design of the instrument, but are instead chosen as part of the survey strategy. We will systematically examine the effects of varying these parameters in Section 7. One of the advantages of intensity mapping with radio telescopes is that substantial fractions of the

sky can be surveyed to a useful depth over the course of only a year or so. This is thanks in part to the relative cheapness of low noise multiple-receiver systems, and the \sim degree-scale primary beams of dishes in most arrays, both of which act to substantially improve survey speed. For much of what follows, we will assume that all experiments can perform between 10 to 25,000 sq. deg. surveys over 10,000 hours total observing time, which is reasonable for a dedicated survey telescope.

We now turn to the effective beam terms. The instrumental resolution in the radial direction is limited by the bandwidth of an individual frequency channel, $\delta\nu$. We approximate the channel bandpass by a Gaussian, which gives an effective beam in the parallel direction,

$$B_{\parallel}(y) = \exp\left(-\frac{(y\delta\nu/\nu_{21})^2}{16 \ln 2}\right);$$

the numerical factor comes from the definition of the full width at half maximum (FWHM), $\sigma = \theta_{\text{FWHM}} \sqrt{8 \ln 2}$. Modern radio receivers can typically be built with narrow channel bandwidths of around 100 kHz or less. Narrow channels allow for more precise removal of artificial radio interference (RFI), for example, but also increase the data rate, which may require expensive increases in correlator performance for interferometers. In practise, the channel bandwidth is not the limiting factor in the radial resolution, as for realistic $\delta\nu$ this is instead determined by the non-linear dispersion scale, σ_{NL} .

The expressions for the dish multiplicity factor, \mathcal{I} , and transverse effective beam, B_{\perp} , depend on whether the array is used as an interferometer or a collection of independent single-dishes, i.e. whether the signals from individual dishes are correlated with one another or not. For *single-dish* experiments that only use the autocorrelation of the signal from each dish, we have

$$\mathcal{I} = \frac{f(\nu)}{N_{\text{b}} N_{\text{d}}}$$

$$B_{\perp}(\mathbf{q}) = \exp\left(-\frac{(q\theta_{\text{B}})^2}{16 \ln 2}\right),$$

where $\theta_{\text{b}} \approx \lambda/D_{\text{dish}}$ is the FWHM of the beam of an individual dish of diameter D_{dish} at some wavelength λ , and N_{d} is the number of dishes in the array. N_{b} is the number of beams, which differs from unity if the experiment has multiple pixels or phased array feeds (PAFs) per dish. Each dish/beam will typically survey a different patch of the sky, thus increasing the survey speed. Any additional frequency dependence of the sensitivity (e.g. due to beam overlap for PAF receivers) can be accounted for by the $f(\nu)$ factor; specific forms of the noise expression appropriate for different types of receiver are given in Appendix D.

For interferometers, in the case where we assume multiple pointings, i.e. $S_{\text{area}} > \text{FOV}$, we have

$$\mathcal{I} B_{\perp}^{-2} = \frac{\text{FOV}}{n(u=q/2\pi)},$$

where $n(u)$ is the number density of samples in the uv plane as a function of $|u|$, and $\text{FOV} \approx (\lambda/D_{\text{dish}})^2$ is the field of view. Each configuration of baselines will lead to a different $n(u)$, although if one assumes constant sam-

pling in uv , it can be approximated by

$$n(u) = \frac{N_d(N_d - 1)}{2\pi(u_{\max}^2 - u_{\min}^2)}, \quad (6)$$

where $u_{\max} = D_{\max}/\lambda$, $u_{\min} = D_{\min}/\lambda$, and D_{\max} , D_{\min} are the lengths of the longest/shortest baselines. The effective beam in the transverse direction is determined by $n(u)$, and so does not need to be defined separately.

For interferometric experiments where the baseline distribution is available, we calculate $n(u)$ specifically for that distribution; otherwise, we use the approximation in Eq. (6). For the former, one has to assume a declination of observation as well as a baseline distribution. As the sky drifts, a set of tracks will be mapped out onto the uv plane. These are typically split into bins of size $(\Delta u)^2 \sim 1/\text{FOV}$, which can be taken to be independent. It is then possible to construct a simple model for $n(u)$ that is good enough for our forecasts (see Appendix C).

2.3. Foreground model

Foreground contamination from the galaxy and extragalactic point sources dwarfs the cosmological HI signal. A number of different methods have been proposed for modelling and subtracting the foregrounds (Oh & Mack 2003; Barkana & Loeb 2005; Santos et al. 2005; Morales et al. 2006; Wang et al. 2006; Gleser et al. 2008; Jelić et al. 2008; Liu et al. 2009; Petrovic & Oh 2011), but in this paper we will simply assume that some sort of method has been applied that removes them, leaving behind some residual contamination whose variance can be modelled as a sum of smooth power spectra. (One could also model instrumental calibration residuals in this way.) Our model for the residual foreground is

$$C^F(\mathbf{q}, y) = \epsilon_{\text{FG}}^2 \sum_X A_X \left(\frac{l_p}{2\pi q} \right)^{n_X} \left(\frac{\nu_p}{\nu_i} \right)^{m_X}. \quad (7)$$

The amplitude and index parameters for four representative foregrounds are given in Table 1, following (Santos et al. 2005).

Subtraction algorithms also introduce a minimum wave number below which cosmological information cannot be extracted. This is because the smooth variation of the foregrounds in frequency is difficult to separate from cosmological modes on scales comparable to the total survey bandwidth, $k_{\text{FG}} \sim 1/(r_\nu \Delta \tilde{\nu}_{\text{tot}})$. A subtraction method that relies on fitting and subtracting smoothly-varying functions will necessarily remove some power from radial cosmological modes larger than this scale as well.

We have introduced an overall scaling, ϵ_{FG} , to parametrise the efficiency of the foreground removal process: $\epsilon_{\text{FG}} = 1$ corresponds to no foreground removal,

Foreground	$A_X [\text{mK}^2]$	n_X	m_X
Extragalactic point sources	57.0	1.1	2.07
Extragalactic free-free	0.014	1.0	2.10
Galactic synchrotron	700	2.4	2.80
Galactic free-free	0.088	3.0	2.15

TABLE 1

FOREGROUND MODEL PARAMETERS AT $\ell_p = 1000$ AND $\nu_p = 130$ MHz, TAKEN FROM SANTOS ET AL. (2005).

while we will probably need $\epsilon_{\text{FG}} \lesssim 10^{-5}$ to be able to extract the cosmological signal. One can also interpret ϵ_{FG} as a measure of how smooth (or correlated) foregrounds are in frequency, and thus how well they can be modelled by smooth deterministic functions. For example, if we assume a Gaussian correlation function along the frequency direction, we have that $\epsilon_{\text{FG}} \sim \exp[-(\Delta\nu/\xi)^2]$, where $\Delta\nu$ is the bandwidth of the redshift bin we are probing and ξ is the correlation length in frequency. We have neglected cross-correlations between different frequencies here, although including them would not be difficult.

This treatment of foreground subtraction is necessarily simplified. For example, we are assuming that the various contributions to the total signal remain uncorrelated, yet all of the subtraction methods proposed so far remove modes that receive contributions from all components. Although the foregrounds will be the dominant part of the subtracted signal by far, this process will nevertheless induce cross-correlations between whatever is left in the residual. Practical experience of foreground subtraction from real data is quite limited so far (see (Chang et al. 2010; Switzer et al. 2013; Masui et al. 2013) for some initial attempts), and so we do not have a good picture of how important various aspects of the foreground problem are yet. For the time being, we believe that our model captures the essential elements of the foregrounds sufficiently well to be of use. We will return to the issue of foreground modelling in Sect. 6.3.

2.4. The Fisher Matrix

We are now in a position to construct the full Fisher matrix. To do this, we need to sum over the number of independent modes in \mathbf{q} and y . If S_{area} is the survey area, FOV is the field of view of a single array element, and $\Delta\tilde{\nu}$ is the (dimensionless) bandwidth in the redshift bin, we have $\Delta y = 2\pi/\Delta\tilde{\nu}$, and $\Delta q = 2\pi/\sqrt{S_{\text{area}}}$ (single-dish) or $\Delta q = 2\pi/\sqrt{\text{FOV}}$ (interferometer). The sum over modes is then given by

$$\int \frac{d^2 q dy}{(\Delta q)^2 \Delta y} \rightarrow U_{\text{bin}} \int \frac{d^2 q dy}{(2\pi)^3},$$

where $U_{\text{bin}} = S_{\text{area}} \times \Delta\tilde{\nu}$. If we define $C^T = C^S + C^N + C^F$, the Fisher matrix for a set of cosmological parameters $\{p_i\}$ is given by

$$F_{ij}^{\text{IM}} = \frac{1}{2} U_{\text{bin}} \int \frac{d^2 q dy}{(2\pi)^3} [\partial_i \ln C^T(\mathbf{q}, y) \partial_j \ln C^T(\mathbf{q}, y)], \quad (8)$$

where the derivatives will only act on C^S , since that is the only term containing parameters of interest.

We can rewrite (8) in terms of physical wave numbers by using the following dictionary: $V_{\text{bin}} = U_{\text{bin}} r_\nu^2$, $\mathbf{q} = \mathbf{k}_\perp r$, and $y = k_\parallel r_\nu$, where $k_\perp = k\sqrt{1-\mu^2}$ and $k_\parallel = k\mu$. We then apply the substitutions

- $U_{\text{bin}} \rightarrow V_{\text{bin}}$
- $\int d^2 q dy \rightarrow 2\pi \int_{-1}^{+1} d\mu \int_{k_{\min}}^{\infty} k^2 dk$
- $q = kr\sqrt{1-\mu^2}$ and $y = kr_\nu\mu$.

We can now express the Fisher matrix in a familiar form to those working on galaxy redshift surveys – a comparison we will pursue in Section 3.

2.5. Experimental configurations

Our focus in this paper is on the lead-up to **Phase I of the SKA**. We consider a portfolio of planned experimental configurations, with the aim of exploring how they will impact constraints on cosmological parameters. Our approach is ecumenical – we try to include as many proposed configurations as possible, although we are limited by what information has been made publicly available.

We will first of all consider three illustrative experimental setups, roughly corresponding to successive ‘generations’ of planned IM experiments. These are:

- **Stage I** – Small, specialised HI experiment focused on a relatively narrow redshift range, intended to provide initial detections of the BAO and other first cosmological results. Stage I experiments are envisaged as either surveys on existing general-purpose arrays, or relatively cheap purpose-built telescopes using multi-feed receivers to improve sensitivity.
- **Stage II** – Larger interferometric experiment with enhanced sensitivity, covering a wider range of redshifts. Stage II experiments are intended to cover a substantial survey volume, with the aim of producing constraints on cosmological parameters that are competitive with contemporary (DETF³ Stage II/III) LSS surveys. They are likely to be either purpose-built HI arrays with a large number of receivers, or surveys on forthcoming ‘SKA-precursor’ arrays such as MeerKAT and ASKAP.
- **Facility** – Survey on a future large array, covering a wide redshift range over most of the sky. Facility-type surveys will compete with other large (DETF Stage IV) experiments to produce the most precise cosmological parameter estimates, and will be able to probe novel HI-only effects for the first time. The only planned experiments of this type so far are the Phase I SKA arrays, although the full CHIME and Tianlai configurations could also fall into this class.

We have chosen representative configurations for each of these classes (see Table 2) that will be used to illustrate the expected progress of HI IM experiments over the next decade. We have also forecasted for the following real (existing, proposed, or plausible) experiments:

- ASKAP**: An SKA pathfinder consisting of thirty-six 12m dishes, each with 36-element PAFs, located at the eventual site of SKA1-SUR in Australia (Johnston et al. 2008).
- BAOBAB**: Proposed compact array of 128 1.6m tiles with 4 dipoles per tile, co-located with GBT or SKA1-MID (Pober et al. 2013).
- BINGO**: A proposed 40m (25m illuminated) multi-receiver single-dish telescope in South America (Battye et al. 2013).
- CHIME**: A proposed array made up of 20×100 m cylinders (20×80 m illuminated), based in British

Columbia, Canada. There is a pathfinder with 2 half-length cylinders, and a planned full experiment with 5 (CHIME Collaboration 2012).

- FAST**: A proposed multi-beam system on the 500m single-dish telescope currently under construction in south-west China (Smoot & Debono 2014).
- GBT**: A 100m single-dish telescope in West Virginia (USA). GBT has already been used for preliminary detections of the HI signal (Chang et al. 2010; Switzer et al. 2013; Masui et al. 2013).
- GBT-HIM**: A planned seven-beam receiver system on GBT (Chang & GBT-HIM Team 2014).
- GMRT**: Array of thirty 45m dishes in Pune, India (Swarup et al. 1991).
- JVLA**: An array of twenty-seven 25m dishes, based in New Mexico, USA (NRAO 2014).
- KAT7**: An SKA pathfinder made up of seven 12m dishes, on the planned site of SKA1-MID (SKA South Africa 2014).
- MeerKAT**: An SKA pathfinder with sixty-four 13.5m dishes, on the site of SKA1-MID (Jonas 2009). Has a choice of two frequency bands.
- MFAA**: A proposal for a mid-frequency aperture array component of Phase II of the SKA (MFAA 2014).
- Parkes**: A single 64m dish (with 13 beams) in NSW, Australia (ATNF 2014).
- SKA1-MID**: A planned SKA Phase I configuration with one hundred and ninety 15m dishes, based in the Northern Cape, South Africa (Dewdney et al. 2013). Can be extended to incorporate the 64 MeerKAT dishes.
- SKA1-SUR**: A planned SKA Phase I configuration with sixty 15m dishes, each with 36-element PAFs, based in Western Australia (Dewdney et al. 2013). Can be extended to incorporate the ASKAP dishes.
- Tianlai**: A proposed array of eight 15×120 m cylinders to be built in north-west China (Chen 2012).
- VLBA**: An array of ten 25m dishes distributed across North America (Napier et al. 1994).
- WSRT + APERTIF**: A proposed upgrade to WSRT that uses a phased array feed (PAF) in the focal plane to produce multiple beams on the sky (Oosterloo et al. 2010).

This is intended to be a relatively exhaustive list of current and planned HI intensity mapping experiments at $z \lesssim 3$, but inevitably some have been omitted due to a lack of publicly-available specifications. We have made our code publicly available,⁴ so forecasts can be performed when specifications become available.

³See Albrecht et al. (2006).

⁴<https://gitlab.com/radio-fisher/bao21cm>

	Experiments	T_{inst} [K]	$N_d \times N_b$	D_{dish} [m]	D_{min} [m]	D_{max} [m]	ν_{crit} [MHz]	$\nu_{\text{max}}^{\text{IM}}$ [MHz]	$\nu_{\text{min}}^{\text{IM}}$ [MHz]	$\Delta\nu^{\text{IM}}$ [MHz]	z_{min}	z_{max}	S_{area} [deg ²]
Ref.	Stage I	50	1×50	30.0	–	–	–	1100	800	300	0.29	0.77	5,000
	• Stage II	35	160×1	4.0	4.0	53.0	1000	1000	600	400	0.42	1.37	2,000
	Facility	20	250×1	15.0	–	–	–	1100	400	700	0.29	2.55	25,000
Existing Facility	GBT	29	1×1	100.0	–	–	–	920	680	240	0.54	1.09	100
	GBT-HIM	33	1×7	100.0	–	–	–	900	700	200	0.58	1.03	1,000
	GMRT	70	30×1	45.0	–	–	–	1420	1000	420	0.00	0.42	1,000
	JVLA	70	27×1	25.0	–	–	–	1420	1000	420	0.00	0.42	1,000
	Parkes	23	1×13	64.0	–	–	–	1420	1155	265	0.00	0.23	5,000
	VLBA	27	10×1	25.0	–	–	–	1420	1200	220	0.00	0.18	5,000
	▲ WSRT + APERTIF	52	14×37	25.0	–	–	–	1300	1000	300	0.09	0.42	25,000
Targeted IM	• BAOBAB-128	40	128×1	1.6	1.6	26.0	–	900	600	300	0.58	1.37	1,000
	BINGO	50	1×50	25.0	–	–	–	1260	960	300	0.13	0.48	5,000
	◇ CHIME	50	1280×1	20.0	–	–	–	800	400	400	0.77	2.55	25,000
	FAST	20	1×20	500.0	–	–	–	1000	400	600	0.42	2.55	2,000
	• MFAA	50	3100×1	2.4	0.1	250.0	950	950	450	500	0.49	2.16	5,000
	◇ Tianlai	50	2048×1	15.0	–	–	–	950	550	400	0.49	1.58	25,000
Pre-SKA	▲ ASKAP	50	36×36	12.0	–	–	–	1000	700	300	0.42	1.03	25,000
	KAT7	30	7×1	13.5	–	–	–	1420	1200	220	0.00	0.18	2,000
	MeerKAT (B1)	29	64×1	13.5	–	–	–	1015	580	435	0.40	1.45	25,000
	MeerKAT (B2)	20	64×1	13.5	–	–	–	1420	900	520	0.00	0.58	25,000
SKA Phase I	SKA1-MID (B1) Autocorr.	28	190×1	15.0	–	–	–	1050	350	700	0.35	3.06	25,000
	◦ SKA1-MID (B1) Interferom.	28	190×1	15.0	–	–	–	1050	350	700	0.35	3.06	100
	SKA1-MID (B2) Autocorr.	20	190×1	15.0	–	–	–	1420	900	520	0.00	0.58	25,000
	◦ SKA1-MID (B2) Interferom.	20	190×1	15.0	–	–	–	1420	900	520	0.00	0.58	50
	▲ SKA1-SUR (B1)	50	60×36	15.0	–	–	710	900	400	500	0.58	2.55	25,000
	▲ SKA1-SUR (B2)	30	96×36	15.0	–	–	1300	1150	650	500	0.23	1.18	25,000
	SKA1-MID + MeerKAT (B1) [†]	–	–	–	–	–	–	1050	350	700	0.35	3.06	25,000
	SKA1-MID + MeerKAT (B2) [†]	–	–	–	–	–	–	1420	900	520	0.00	0.58	25,000

TABLE 2

TELESCOPE CONFIGURATIONS USED IN THIS PAPER. THE ASSUMED OBSERVING MODE OF EACH TELESCOPE IS DENOTED BY: () SINGLE-DISH; (▲) SINGLE-DISH WITH PHASED ARRAY FEED; (◦) DISH INTERFEROMETER; (•) DENSE APERTURE ARRAY; (◇) CYLINDER INTERFEROMETER. SOME INSTRUMENTS CAN OPERATE OVER A WIDER FREQUENCY RANGE THAN SHOWN HERE; WHERE THIS IS THE CASE, OUR VALUES CORRESPOND TO THE MOST APPROPRIATE ν_{max} FOR IM, OR WE INCLUDE MULTIPLE BANDS. [†]FOR COMBINED ARRAYS, IN REDSHIFT BINS WHERE THE BANDS OVERLAP, WE FIND T_{inst} , D_{dish} BY AVERAGING THE VALUES FOR EACH SUB-ARRAY, WEIGHTED BY THE NO. OF DISHES.

The instrumental parameters used for each experiment are listed in Table 2. The survey area is chosen (between 10 and 25,000 deg²) to maximise the DE figure of merit (see Sect. 5), and the survey time is assumed to be $t_{\text{tot}} = 10^4$ hours for all experiments. For interferometers, we use either the baseline density calculated from the actual array layout (see App. C), or the $n(u) \sim \text{const.}$ approximation of Eq. (6). For simplicity, we consider only single-dish mode for some telescope arrays, even though they are capable of interferometric measurements.

It is useful to be able to compare the performance of IM experiments with competing probes, such as galaxy redshift surveys. To this end, we also produce forecasts for a fiducial DETF Stage IV spectroscopic galaxy redshift survey, similar to DESI, Euclid or WFIRST, with an expected yield of $\sim 6 \times 10^7$ spectroscopic redshifts. We take the survey to cover roughly a third of the sky ($f_{\text{sky}} = 0.35$) over a redshift range of $0.65 \leq z \leq 2.05$, with redshift distribution taken from Amendola et al. (2013) (Euclid reference case, Table 1.3). The bias is taken to evolve as $b(z) = \sqrt{1+z}$. We forecast for the same set of cosmological parameters as IM experiments, with the same fiducial values, including relevant nuisance parameters like σ_{NL} . We use the redshift scaling from Smith et al. (2003) to set $k_{\text{max}} = k_{\text{NL}}(1+z)^{2/(2+n_s)}$, and

choose $k_{\text{min}} = 2\pi/(V_{\text{bin}})^{\frac{1}{3}}$. Constraints are quite sensitive to the choice of k_{max} (White et al. 2008), but mostly insensitive to k_{min} (if chosen sufficiently small).

2.6. Prior information

Intensity mapping experiments cannot constrain all cosmological quantities of interest on their own; information from other probes must be added in order to break degeneracies and improve precision. High quality data are already available from a range of other sources, including CMB temperature and polarisation anisotropies at high and low- ℓ , galaxy redshift surveys at high and low redshift, weak gravitational lensing, and supernova distance measurements. By the time of the first intensity mapping surveys with the SKA, however, the number of experiments for each of these probes, and their precision, will have risen sharply.

While the best constraints will ultimately be obtained by combining information from all relevant experiments, our intention here is not to provide an exhaustive account of the expected state of observational cosmology in ten years' time. Instead, we will (conservatively) focus only on the CMB as an external probe in this paper.

CMB data provides a high-redshift distance measurement that is vital for anchoring the low- z distance mea-




FIG. 1.— Schematic illustration of the ranges of radial and transverse wavenumbers that the two types of experiment are sensitive to. The dotted lines show the ranges in absolute wavenumber $|k|$; single-dish experiments are sensitive to smaller $|k|$ due to their lower k_{\perp}^{min} , while interferometers can see larger $|k|$ on account of their high angular resolution. The two types of experiment are complementary in terms of their angular sensitivity, but are subject to the same constraints in frequency space.

sures that most effectively probe dark energy. It also yields information about the shape and normalisation of the matter power spectrum, the matter content at $z \simeq 1090$, and the physical scale of the BAO.

We use the DETF Planck prior from Albrecht et al. (2009), which assumes temperature and E-mode polarisation data over 70% of the sky for the 70, 100, and 143 GHz channels out to $\ell = 2000$. We rescale the Fisher matrix to reflect our different fiducial cosmology, and then project it out to the following variables (fixing all others):

$$\{h, \Omega_b h^2, \Omega_{\text{DE}}, \Omega_K, w_0, w_a, n_s, \sigma_8\}.$$

Note that the optical depth to last scattering, τ , which depends on the cosmic reionisation history, has been marginalised over in the original DETF Fisher matrix. We ignore constraints from CMB lensing, high- ℓ CMB experiments and B-mode polarisation, although in principle these would provide additional information on h , σ_8 , n_s , and the effective number of neutrino species.

3. COMPARISON WITH GALAXY REDSHIFT SURVEYS

It is instructive to compare IM surveys of large scale structure with ‘conventional’ redshift surveys, by which we mean surveys that catalogue individual galaxies in angle and redshift. These offer some of the most stringent constraints on cosmological parameters to date, and are likely to do so for some time as experiments like the Dark Energy Survey (DES) (Frieman & Dark Energy Survey Collaboration 2013) and Euclid (Amendola et al. 2013) come online.

In this section, we will compare IM and galaxy redshift surveys by looking directly at constraints on the power spectrum, $P(k)$. To do this we divide up a range of wavenumbers into bins, $\Delta_a = [k_a, k_{a+1}]$, and assign a constant value, P_a to the power spectrum in each bin. The exercise is then to forecast errors for each P_a .

The main quantities that describe a redshift survey are the survey area, S_{area} , and the number density of sources




FIG. 2.— The normalised effective volume $V_{\text{eff}}(k_{\perp}, k_{\parallel})/V_{\text{bin}}$ at $z \approx 1$, for SKA1-MID Band 1 in single-dish mode (black contours) and interferometer mode (shaded blue contours). Foregrounds have not been included, but the effective minimum k_{\parallel} , given by $k_{\text{FG}} \simeq k_{\text{BW}}$, is shown as a dashed grey line. The contours are for values $[0.9, 0.5, 0.1, 0.01, 0.001]$, where 1.0 is the maximum (implying a cosmic variance-limited measurement). Only the last three contours (i.e. < 0.5) appear for the interferometer mode due to its lower sensitivity.

as a function of redshift, $\bar{n}(z)$, which in turn allows us to define a survey depth. As with the IM survey, we can define the survey volume, $V_{\text{bin}} \simeq r^2(\bar{z})(dr/dz)\Delta z S_{\text{area}}$ (where \bar{z} is the mean redshift of the survey). The Fisher matrix is then (Feldman et al. 1994; Tegmark et al. 1998)

$$F_{ij}^{\text{gal}} = \frac{1}{2} \int_{k_{\text{min}}}^{k_{\text{max}}} \frac{d^3 k}{(2\pi)^3} [\partial_i \ln P_{\text{tot}}(\mathbf{k}) \partial_j \ln P_{\text{tot}}(\mathbf{k})] V_{\text{eff}}(\mathbf{k})$$

$$V_{\text{eff}} = V_{\text{bin}} \left[\frac{\bar{n}(z) P_{\text{tot}}(\mathbf{k})}{1 + \bar{n}(z) P_{\text{tot}}(\mathbf{k})} \right]^2. \quad (9)$$

Shot noise plays a crucial role, dominating if \bar{n} is too small; as does cosmic variance, via the V_{bin} term. The effective volume tells us how well different parts of Fourier space are sampled. Applying (9) to the binned power spectrum, we recover the well-known result

$$\left\langle \left(\frac{\Delta P_a}{P_a} \right)^2 \right\rangle = \left[\frac{1}{2} \int_{V_n} \frac{d^3 k}{(2\pi)^3} V_{\text{eff}}(\mathbf{k}) \right]^{-1}, \quad (10)$$

which has V_{eff} at the heart of the expression. It is only in regions where $\bar{n}P \gg 1$ that the power spectrum can be measured well. In this regime the fundamental limitation becomes cosmic variance, which is set by the number of modes sampled in each bin, $N_a \sim V_{\text{bin}} k_a^2 \Delta k_a / 2\pi^2$, where Δk_a is the width of the corresponding bin.

The equivalent expression for an IM survey is

$$\left(\frac{\Delta P_a}{P_a} \right)^2 = \left[\frac{1}{2} U_{\text{bin}} \int_{V_n} \frac{d^2 q dy}{(2\pi)^3} \left(\frac{C^S(\mathbf{q}, y)}{C^T(\mathbf{q}, y)} \right)^2 \right]^{-1} \quad (11)$$

By analogy with (9) and (10), we can define an effective volume, $V_{\text{eff}}^{\text{IM}}$, and a pseudo number density, $\bar{n}^{\text{IM}} P(\mathbf{k}) \equiv C^S/C^N$, such that

$$\bar{n}^{\text{IM}}(z, \mathbf{k}) = \left(\frac{T_b}{T_{\text{sys}}} \right)^2 \frac{t_{\text{tot}} \Delta \nu}{U_{\text{bin}}} B_{\perp}^2 B_{\parallel} \mathcal{I}^{-1}. \quad (12)$$

The foreground term has been left out for the time being. With (12) in hand, we are now in a position to exploit the analogy with conventional redshift surveys to better understand the properties of IM experiments. An illustration of the various scales relevant to single-dish and interferometric experiments is shown in Fig. 1, and an example of $V_{\text{eff}}^{\text{IM}}$ is shown in Fig. 2.

The first thing to notice is that $V_{\text{eff}}^{\text{IM}}$ is highly anisotropic. In the case of a single dish, information on angular scales smaller than the instrumental beam is washed out; for dishes of diameter D_{dish} , scales $k_{\perp} \gtrsim D_{\text{dish}}/r\lambda$ are suppressed. For an interferometer, it is possible to probe much smaller angular scales (up to $k_{\perp} \sim 2\pi u_{\text{max}}/r$), although the transverse Fourier plane will be sampled much less homogeneously than for a single dish, depending on the instrument's uv coverage.

Along the radial direction, we expect foreground subtraction to throw away information on scales of order the total bandwidth, $k_{\parallel} \lesssim k_{\text{FG}}$, and on smaller scales non-linear velocities smear out information for $k \gtrsim 0.15 \text{ Mpc}^{-1}$. As discussed in Section 2.2, the channel bandwidth also imposes an effective radial beam, although this is generally at higher k_{\parallel} than the non-linear cutoff.

Using Eq. (11), we can now piece together how well the power spectrum can be measured by various experiments. On large scales, cosmic variance and the survey size are the limiting factor. For a single-dish experiment with a relatively isotropic survey volume, we can sample the longest wavelength modes in the radial and transverse directions equally well, so that $(\Delta P/P)^2 \propto k^{-3}$. On smaller scales, the beam size will severely limit how small a transverse scale we can probe, so only radial modes will be properly sampled, implying $(\Delta P/P)^2 \propto k^{-1}$. In the radial direction, we will eventually come up against the non-linear velocity scale, preventing us from extracting information on scales smaller than σ_{NL} .

For an interferometric experiment, the situation is reversed, and is in some sense complementary, as shown in Fig. 1. With sufficiently long baselines, it is possible to probe very small angular scales. We then expect to have $(\Delta P/P)^2 \propto k^{-3}$ up until the non-linear scale is reached in the radial direction, beyond which only transverse scales contribute, such that $(\Delta P/P)^2 \propto k^{-2}$ until the maximum transverse resolution is hit. Conversely, interferometers are fundamentally unable to probe any scale larger than that corresponding to their minimum baseline (which for a dense array is roughly the dish diameter, which gives the beam size in single-dish mode).

In summary, the important scales for IM are:

$$\begin{aligned} k_{\parallel}^{\text{min}} &\sim k_{\text{FG}} = 2\pi/(r_{\nu}\Delta\tilde{\nu}_{\text{tot}}) \\ k_{\parallel}^{\text{max}} &\sim k_{\text{NL}} = 1/\sigma_{\text{NL}} \\ k_{\perp}^{\text{min}} &\sim k_{\text{area}} = 2\pi/\sqrt{r^2 S_{\text{area}}} \quad (\text{single dish}) \\ &\sim k_{D_{\text{min}}} = 2\pi D_{\text{min}}/r\lambda \quad (\text{interferom.}) \\ k_{\perp}^{\text{max}} &\sim k_{\text{FOV}} = 2\pi D_{\text{dish}}/r\lambda \quad (\text{single dish}) \\ &\sim k_{D_{\text{max}}} = 2\pi D_{\text{max}}/r\lambda \quad (\text{interferom.}) \end{aligned}$$

The redshift dependence of the transverse scales for an example setup is shown in Fig. 3. Note that these are only the minimum/maximum scales that can be probed




FIG. 3.— Redshift evolution of the minimum/maximum transverse scales (filled regions) for illustrative interferometer (blue) and single-dish (red) experiments. The BAO are plotted for comparison. The dishes have diameter $D_{\text{dish}} = 15\text{m}$, the min./max. interferometer baselines are $D_{\text{min}} = 15\text{m}$ and $D_{\text{max}} = 1000\text{m}$, and the survey has bandwidth $\Delta\nu = 600 \text{ MHz}$ and area $S_{\text{area}} = 25,000 \text{ sq. deg.}$. The shaded grey region denotes superhorizon scales, $k < k_H = 2\pi/r_H$.

in principle by a given instrument – the sensitivity to scales within these ranges will vary, so it may not be possible to constrain the more extreme scales in practise.

4. EXPANSION, GROWTH, AND THE ACOUSTIC PEAK

We begin our exploration of the capabilities of IM experiments by focusing on a few key observables. These variously constrain the growth of large-scale structure and the expansion and geometry of the Universe: The positions of the acoustic peaks and the overall shape of the power spectrum in both the radial and transverse directions can be used as distance indicators to place constraints on $D_A(z)$ and $H(z)$, and redshift space distortions make it possible to measure the growth rate, $f(z)$.

In addition to being of intrinsic cosmological interest, these also serve as useful models for other observables. For example, the detection of the acoustic peaks is a comparable problem to measuring other ‘shape’ features of the power spectrum, such as scale-dependent bias. We will therefore devote this section to understanding the detailed characteristics of the measurements on these observables that can be made with IM experiments. Throughout, we will forecast for the following parameter set (without any external priors):

$$\{A(z), [b_{\text{H}}\sigma_8](z), [f\sigma_8](z), D_A(z), H(z), \sigma_{\text{NL}}\}.$$

4.1. Detectability of Baryon Acoustic Oscillations

The Baryon Acoustic Oscillations are a ‘statistical standard ruler’ that forms the primary distance measure in surveys of large-scale structure. We can get an idea of the detectability of the BAOs by looking at the fractional errors on $P(k)$, using Eq. (11). These are shown for the reference surveys in Fig. 4, and are overplotted on the BAO wiggle function (to be defined shortly) in Fig.




FIG. 4.— Fractional constraints on $P(k)$ for the set of reference experiments, combined over the whole redshift range of each experiment, with 20 bins per decade in k .

5. All three IM surveys are capable of strongly detecting the BAO feature when the constraints are combined over their full redshift ranges. Facility approaches the cosmic variance limit (represented by the DETF Stage IV survey out to $k \sim 0.1 \text{ Mpc}^{-1}$) over a substantial fraction of the scales relevant to the BAO, mostly due to the sensitivity of its single-dish component. This also helps to put sub-10% level constraints on the power spectrum on scales slightly larger than the matter-radiation equality peak, $k_{\text{eq}} \approx 10^{-2} \text{ Mpc}^{-1}$. Its interferometric component provides constraints on smaller scales, achieving $\sim 10\%$ errors on $P(k)$ out to $k \approx 1 \text{ Mpc}^{-1}$.

The interferometric Stage II survey is sensitive to generally smaller scales, but still achieves good constraints on the BAO thanks to its coverage out to intermediate redshifts ($z \sim 1.4$). The Stage I survey can comfortably detect the BAO despite its significantly lower sensitivity than Facility, but leaves smaller scales unconstrained.

Alternatively, one can look at the detectability of the BAO feature as a whole. We follow a similar approach to (Blake & Glazebrook 2003) and split the matter power spectrum, $P(k)$, into a ‘smooth’ part, $P_{\text{smooth}}(k)$, and an oscillatory part,

$$f_{\text{bao}}(k) = \frac{P(k) - P_{\text{smooth}}(k)}{P_{\text{smooth}}(k)}. \quad (13)$$

We then introduce an amplitude parameter, A , such that

$$P(k) = [1 + A f_{\text{bao}}(k)] P_{\text{smooth}}(k). \quad (14)$$

Constraints on A therefore give a measure of the detectability of the BAO feature.

The splitting of $P(k)$ between smooth and oscillatory parts is somewhat arbitrary. We attempt to construct a ‘purely oscillatory’ $f_{\text{bao}}(k)$ – i.e. one that lacks a smooth overall trend in k – as follows. First, we use CAMB to calculate $P(k)$ for the fiducial cosmological model over a range of sample points in k . We then choose two reference values of k that bound the region in which the oscillations are significant ($k \approx 0.02$ and 0.45 Mpc^{-1} for our fiducial cosmology), and construct a cubic spline for $\log P(k)$ as a function of $\log k$ using all points *outside* that region. Next, we construct a preliminary os-




FIG. 5.— Forecast constraints on the BAO wiggles, combined over the whole redshift range for each of the reference surveys.

cillatory function by dividing the sampled $P(k)$ by the splined function (not its logarithm), then fit another cubic spline to the result and find the zeros of its second derivative with respect to k . These are the points at which the first derivatives of the oscillatory function are maximal/minimal, and in some sense define ‘mid-points’ of the function – its overall trend. We construct a cubic spline through these too, and then divide the preliminary oscillatory function by it to ‘de-trend’. This leaves $f_{\text{bao}}(k)$ as the final result (Fig. 5). Unlike other methods, which look at ratios of the form $P(k, \Omega_b \neq 0)/P(k, \Omega_b = 0)$ to pick out oscillations (Rassat et al. 2008), this method is essentially model-independent for a given fiducial $P(k)$.

The constraint on the overall amplitude of the BAO feature, A , is plotted as a function of redshift for the reference surveys in Fig. 6. Facility is capable of $> 3\sigma$ detections of the BAO feature out to $z \approx 1.5$, but makes progressively weaker detections at higher redshift, predominantly due to its limited angular resolution in single-dish mode. In comparison, the Stage II survey’s constraints degrade much less rapidly with redshift, owing to its greater sensitivity to smaller angular scales (which translate to intermediate physical scales at higher z).

Fig. 7 plots the errors on $P(k)$ for Facility as a function of both scale and redshift. For $k \gtrsim 0.1 \text{ Mpc}^{-1}$, most of the information comes from low redshifts, where the amplitude of the power spectrum is largest. At smaller k , however, the volume of the redshift bin begins to matter, as the increase in bin volume with z allows progressively larger scales to be probed. For Facility, the constraints




FIG. 6.— Fractional errors on $A(z)$, $f\sigma_8(z)$, $D_A(z)$, and $H(z)$, as a function of redshift.

on intermediate BAO scales ($k \sim 0.07 \text{ Mpc}^{-1}$) come from a mixture of low and intermediate redshift bins, with the high redshift bins taking over on larger scales.

4.2. Constraints on $D_A(z)$ and $H(z)$

The angular diameter distance, $D_A(z)$, and the expansion rate, $H(z)$, are measures of distance in the transverse and radial directions respectively. To include them in our forecasts, we introduce (Blake & Glazebrook 2003)

$$\alpha_{\perp} \equiv \frac{r^{\text{fid}}}{r} = \frac{D_A^{\text{fid}}(z)}{D_A(z)} \quad (15)$$

$$\alpha_{\parallel} \equiv \frac{r_{\nu}^{\text{fid}}}{r_{\nu}} = \frac{H(z)}{H^{\text{fid}}(z)} \quad (16)$$

where r^{fid} and r_{ν}^{fid} are the fiducial ΛCDM values of r and r_{ν} at a given redshift. We then replace $q \rightarrow \alpha_{\perp}q$ and $y \rightarrow \alpha_{\parallel}y$ in Eq. (4) to get

$$C^S(\mathbf{q}, y) = T_b^2 \frac{\alpha_{\perp}^2 \alpha_{\parallel}}{r^2 r_{\nu}} F_{\text{RSD}} \left(\frac{\alpha_{\perp} \mathbf{q}}{r}, \frac{\alpha_{\parallel} y}{r_{\nu}} \right) D^2(z) \times P \left(k = \sqrt{\left(\frac{\alpha_{\perp} \mathbf{q}}{r} \right)^2 + \left(\frac{\alpha_{\parallel} y}{r_{\nu}} \right)^2} \right). \quad (17)$$

The distance measures enter this expression in four places: (a) an overall factor of $\alpha_{\perp}^2 \alpha_{\parallel}$, related to the physical volume of the survey; (b) a distortion of the angular (μ) dependence of the RSDs; (c) a shift in the non-linear

cutoff scale of the RSDs; and (d) a shift of the whole isotropic power spectrum, $P(k)$, that can be further subdivided into corresponding shifts in the BAO feature and smooth power spectrum. The latter two, (c) and (d), are both due to the remapping of k , shown explicitly in the argument of $P(k)$ in Eq. (17) (Samushia et al. 2011).

Due to modelling uncertainties and degeneracies with other parameters, it is not necessarily desirable to use all of these terms to measure distances from real data. The BAO feature is the standard choice of distance measure, owing to its robustness to systematic error; the acoustic scale shifts only slightly when non-linearities are introduced (Smith et al. 2008; Crocce & Scoccimarro 2008), and smooth variations such as scale-dependent bias also have a relatively minor impact (Zhang 2008) (although corrections must still be made for precision measurements). Anisotropies of the correlation function can also be used to measure distances (Alcock & Paczynski 1979; Kaiser 1987), through redshift space distortions (b) and the Alcock-Paczynski effect (a, c, and d), although these are more sensitive to the detailed modelling of the power spectrum (Reid et al. 2012). Thus, a particularly conservative analysis might only derive distances from the BAO, and discard information from the other terms.

Fig. 8 shows the effect of neglecting some of the distance terms for the Facility experiment. By using only the BAO, one is discarding a substantial amount of useful information, as shown by the comparatively poor constraints on D_A and H . This is partially compensated by the reduced risk of systematic error, and the improved growth rate constraints that are due to weaker degenera-




FIG. 7.— Fractional constraints on $P(k)$ in each redshift bin, for the Facility experiment. The thick black line is the total constraint, summed over all redshift bins.

cies with other parameters when compared to the other distance terms. Including broadband information – i.e. the shift in (smooth) $P(k)$ – reduces the error on D_A by a factor of 2 – 5 over the entire redshift range, and is especially beneficial at higher z , where the BAO-only constraints degrade rapidly due to the limited angular sensitivity of the telescope. Adding the RSD terms helps to distinguish between the radial and transverse directions, which also reduces the error on $H(z)$.

Disregarding any of the distance terms can also substantially alter the correlation structure of the Fisher matrix, which has a knock-on effect on constraints for other parameters. This can be seen clearly for $f(z)$ in Fig. 8, which has a substantial scatter in fractional error depending on which distance measures are used. To most accurately reflect the interdependencies of the various cosmological parameters, we will use all of the distance measure terms in what follows. The increased uncertainty that comes from marginalising over nuisance parameters such as σ_{NL} helps compensate for the increased risk of systematic error that attends some of the measures, which means that this is not too optimistic of us.

Fig. 6 shows the fractional constraints that can be achieved on $D_A(z)$ and $H(z)$ for the full set of reference surveys. Facility measures the expansion rate to better than 1% out to $z \approx 1.7$, and stays within 2% as far out as $z = 2.5$. This is roughly a factor of two worse than the DETF Stage IV reference survey, although the redshift range covered by Facility is significantly larger. Stage II also obtains $\sim 2\%$ constraints across its full redshift range, and Stage I hovers around 3%.

The picture is somewhat different for $D_A(z)$. The fractional errors for Facility increase from $\sim 1.5\%$ at $z \approx 0.4$ up to 4% at $z = 2$, compared with the relatively flat errors for the galaxy survey that remain below 1% for most of the redshift range. Stage II’s errors are also relatively flat as a function of z , staying at around the 2.5% mark.

The limited angular resolution of the single-dish experiments is the cause of this behaviour. $H(z)$ is most sensitive to the resolution in the radial (frequency) direction, which is essentially the same for all experiments and does not evolve appreciably with redshift (being set




FIG. 8.— Fractional errors on $H(z)$, $D_A(z)$, and $f\sigma_8(z)$ for the Facility experiment, for various combinations of distance terms being switched on and off in the Fisher matrix calculation. The RSDs are taken to include both the angle-dependence and non-linear cut-off terms (see text).

by the non-linear scale rather than the channel bandwidth, as discussed in Sect. 3). The $D_A(z)$ constraints depend more on the sensitivity to transverse physical scales, however, which differs between interferometers and single-dish experiments. For single-dish, k_{\perp}^{max} tends to be relatively small even at $z = 0$ for moderately-sized dishes, and continues to decrease (shift to larger scales) as z increases, as shown in Fig. 3. As this happens, useful distance information from smaller scales is lost. The same happens for interferometers, but k_{\perp}^{max} is typically much larger, so the most useful transverse scales remain resolved. In fact, since k_{\perp}^{min} is also decreasing, additional distance information becomes available from larger scales.

It is worth noting that the effect would be different if only the BAO were being used as distance measures – both the $D_A(z)$ and $H(z)$ constraints would be affected by the loss of resolution in the transverse direction (Sánchez et al. 2013). To see this, consider a simplified model of the correlation function consisting of the sum of a smooth component and a feature, $\xi_{\text{BAO}}(r) = A \exp[-(r - r_{\text{BAO}})^2/2\sigma^2]$. The response to the loss of resolution along the transverse direction corresponds to convolving the correlation function with a window function such that

$$\tilde{\xi}_{\text{BAO}}(r_{\parallel}, \mathbf{r}_{\perp}) = \int d^2 r'_{\perp} W(\mathbf{r}_{\perp} - \mathbf{r}'_{\perp}) \xi_{\text{BAO}}(\sqrt{r_{\parallel}^2 + r_{\perp}^2})$$




FIG. 9.— Fractional errors on the volume distance, $D_V(z)$, and Alcock-Paczynski distortion, $F(z)$, for our reference surveys.

which can be rearranged to have the form

$$\tilde{\xi}_{\text{BAO}}(r_{\parallel}, \mathbf{r}_{\perp}) = \int d^2\Delta W(\Delta) \xi_{\text{BAO}} \left(\sqrt{r^2 + 2\mathbf{\Delta} \cdot \mathbf{r}_{\perp} + \Delta^2} \right).$$

The smoothed correlation function remains a function of r and r_{\perp} , which means that the degradation of signal due to the smoothing is almost democratically taken up by both the transverse and parallel directions. This can be seen explicitly in Fig. 8 for the BAO-only curves for $D_A(z)$ and $H(z)$, which both show a much stronger evolution with redshift than any other combination of distance measures. The inclusion of additional distance measures is therefore necessary to help limit the effect of the poor angular resolution of single-dish experiments.

The various distance measures for LSS surveys depend on combinations of D_A and H rather than constraining them individually. For example, moments of the correlation function (Reid et al. 2012) give the volume distance and Alcock-Paczynski terms,

$$D_V(z) = \left[(1+z)^2 D_A^2 \frac{cz}{H(z)} \right]^{\frac{1}{3}} \quad (18)$$

$$F(z) = (1+z) D_A(z) H(z) / c. \quad (19)$$

In some sense, these quantities define redshift-dependent figures of merit – many other studies forecast directly in terms of D_V , for example, and surveys are compared in terms of the errors that they can achieve on this parameter at a given redshift. We have therefore presented results for both D_A and H (Fig. 6), and D_V and F (Fig. 9) to facilitate comparison with previous studies.

4.3. Constraints on the growth rate, $f(z)$

The other key observable provided by LSS surveys is the growth rate, which can be measured from the anisotropy of the correlation function in redshift space.

The growth rate has two major roles: first, as another measure of distance, since it can be related to the expansion rate (albeit in a model-dependent way); and second,

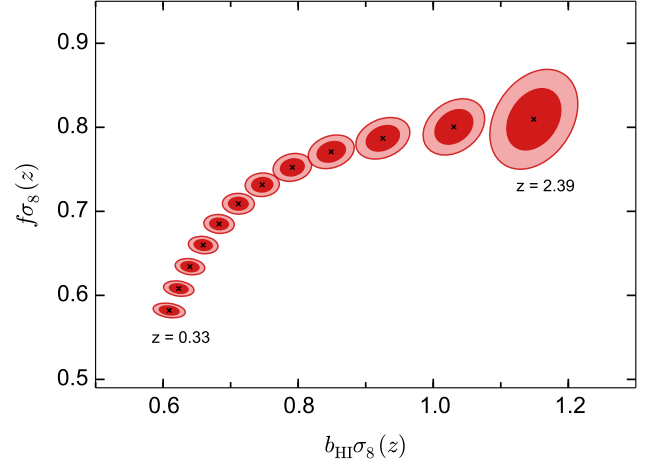


FIG. 10.— Constraints on $b_{\text{HI}}\sigma_8$ and $f\sigma_8$ as a function of redshift, for the Facility survey.

as a non-geometric probe of gravity over cosmic time. The former is useful for helping to break parameter degeneracies that can crop up when only geometric distance measures are used. The latter is of importance in distinguishing theories of dark energy and modified gravity (Hall et al. 2013; Masui et al. 2010); it is often the case that theories can be made to have the same expansion history, for example, but differ in growth rate. (We will return to this in Sect. 5.4.)

In what follows, we will concentrate on the linear growth rate, $f(z)$. HI intensity mapping experiments are also capable of probing non-linear growth, but modelling uncertainties on small scales reduces their usefulness for constraining cosmological parameters. Non-linear effects are discussed in more detail in Sects. 5.4 and 6.2.

The linear growth rate appears in two places in Eq. 17; explicitly, in the angle-dependent RSD factor, and implicitly, through the linear growth factor, $D(z)$, that gives the redshift evolution of the power spectrum. The two are related by $f(z) = d \log D / d \log a$. Because $D(z)$ appears as an overall factor of the signal covariance, it is degenerate with other parameters, making it hard to measure in an unbiased way. We will assume that $D(z)$ provides no new information on the growth rate here (i.e. we neglect its derivative w.r.t. f in the Fisher matrix).

Various other degeneracies crop up when measuring the growth rate from RSDs. The signal covariance is proportional to $C^S \propto [b_{\text{HI}}(z) + f(z)\mu^2]^2 D^2(z) T_b^2(z) \sigma_8^2$. If no functional form is assumed for any of these terms (i.e. they are left free in each redshift bin), there are only two quantities that can be uniquely distinguished from this expression: an overall amplitude, and a factor of the angle dependence, e.g. $C^S \propto \kappa^2(z) [1 + \beta(z)\mu^2]^2$, where $\beta = f(z)/b_{\text{HI}}(z)$ and $\kappa = \sigma_8 D(z) T_b(z) b_{\text{HI}}(z)$. Alternatively, one can merge the linear growth factor with the overall normalisation to give a redshift-dependent normalisation, $\sigma_8(z) = \sigma_8 D(z)$, and write the two RSD functions as $T_b f \sigma_8$ and $T_b b_{\text{HI}} \sigma_8$.

Either way, there are at least three unknowns to be determined from two functions, so it is clear that more information is needed to unpick the degeneracy. The CMB gives a prior on $\sigma_8(z \simeq 1090)$, $D(z)$ can in principle be determined from $f(z)$, and several models for the

bias and brightness temperature exist, although there is significant disagreement between them (see Sect. 6.1). The brightness temperature will be measurable from the non-fluctuating part of the HI signal when future IM experiments come online, though, so this can reasonably be taken as a given quantity – $T_b(z)$ is fixed to its fiducial form throughout this paper. We resolve the remaining degeneracy by treating the combinations $f(z)\sigma_8(z)$ and $b_{\text{HI}}(z)\sigma_8(z)$ as independent parameters, with both being free functions of redshift. Fig. 10 shows the joint constraints on them as a function of z for the Facility survey.

Fig. 6 shows the constraints on $f\sigma_8(z)$ that can be achieved by our set of reference surveys. Sub-2% errors are possible for the Facility and Stage II experiments out to a redshift of $z \sim 1.2$, despite (pessimistically) taking the bias to be a completely free function of redshift. As shown in Fig. 8, constraints on the growth rate are sensitive to the choice of other distance measures; for Facility at least, using only BAO to measure distances would result in a $\sim 50\%$ reduction in the error on $f\sigma_8(z)$ across the whole redshift range, albeit at the cost of significantly degraded $H(z)$ and $D_A(z)$ measurements.

As shown in Fig. 6, the errors on $f\sigma_8(z)$ for the Stage I and Facility surveys increase significantly with redshift, while the evolution is less severe for Stage II. As with the angular diameter distance (see previous section), this is mostly due to the limited angular resolution of the dish-based surveys.

5. COSMOLOGICAL PARAMETERS

In the previous section, we assessed how well HI intensity mapping experiments will be able to measure the geometry, expansion and growth rate of the Universe. We will now discuss how these map to constraints on the cosmological parameters that characterise the standard Λ CDM model, including extensions such as a time-varying equation of state of dark energy, non-zero spatial curvature, and a modified growth index.

One can map the functions of redshift into the set of cosmological parameters using a simple linear transformation of the Fisher matrix,

$$F(\beta) = \sum_i M_i^T F_i(\alpha) M_i, \quad (20)$$

where $\alpha = \{f(z), D_A(z), H(z)\}$ are the old parameters, $\beta = \{h, \Omega_{\text{DE}}, \Omega_K, w_0, w_a, \gamma\}$ are the new parameters, F_i is the Fisher matrix in a bin with redshift z_i , and the transformation matrix is given by $M_{jk}(z_i) = \partial\alpha_j(z_i)/\partial\beta_k$ (Albrecht et al. 2009). The derivatives required for the transformation matrix are all analytical; for completeness, we present them in Appendix E.

To complete the set of cosmological parameters, we must also include information on the shape and normalisation of the initial power spectrum, $\{n_s, \sigma_8\}$. These parameters are derived directly from the signal model of Eq. (4), and do not depend on the functions of redshift from the previous section (i.e. we have separated $f(z)$ and σ_8). Note that we do not use the shape of the power spectrum to constrain any other parameters, such as Ω_M or Ω_b , even though in principle it does depend on them.

Carrying over the remaining parameters from the previous section, the full set is now

$$\{h, \Omega_{\text{DE}}, \Omega_K, w_0, w_a, n_s, \sigma_8, \gamma, A(z), b_{\text{HI}}(z), \sigma_{\text{NL}}, \omega_b\}.$$

The baryon density, $\omega_b = \Omega_b h^2$, is not constrained directly by HI experiments, but is included in the Planck prior. The total matter density (CDM + baryons) is fixed by $\Omega_M = 1 - \Omega_K - \Omega_{\text{DE}}$, so we do not include it separately. The HI bias is free in each redshift bin, and we have taken σ_8 to be constant in redshift.

In what follows, we focus on the higher-end reference experiments, Stage II and Facility, although marginal (1D) constraints are provided for all of the experiments listed in Sect. 2.5. We will also consider the effect of adding prior information from the CMB.

5.1. ‘Vanilla’ Λ CDM

The current consensus is that cosmological data are well-described by the a flat Λ CDM model of structure formation that can be characterised in terms of six parameters: the Hubble parameter, $H_0 = 100h \text{ km s}^{-1} \text{ Mpc}^{-1}$, the density of dark energy (or cosmological constant), Ω_{DE} , the physical density of baryons, ω_b , the linear amplitude of density fluctuations, parametrised by σ_8 , the spectral index of primordial density perturbations, n_s , and the optical depth to last scattering, τ . In this section, we examine the constraints that intensity mapping experiments will be able to put on this model when combined with CMB data from Planck. Parameters that extend the ‘vanilla’ Λ CDM model ($w_0, w_a, \Omega_K, \gamma$) are fixed to their fiducial values in this section, and $\Omega_{\text{DE}} = \Omega_\Lambda$.

Fig. 11 presents forecasts for five of the six parameters for the Facility experiment, compared with Planck-only and the DETF Stage IV galaxy redshift survey. Although the reionisation history will have a significant role in the evolution of the HI density and bias, we are focusing on sufficiently late times that our constraints will essentially be insensitive to variations of τ , within current constraints, and so we leave it out of the plot. (In fact, it has already been marginalised over in the Planck prior Fisher matrix.) We do not directly constrain ω_b with IM experiments either, but as it is strongly correlated with other parameters in the Planck prior, we leave it in.

As expected, there is a modest improvement over Planck alone by a factor of a few. The Planck-only constraints are mostly limited by strong correlations between parameters, so the role of the IM survey is primarily to break degeneracies. Future high resolution experiments such as ACTPol (Niemack et al. 2010) and SPTpol (Austermann et al. 2012) will be able to measure weak lensing of the CMB to sufficient accuracy that constraints from the CMB *alone* should be competitive

Experiments	h / 10^{-3}	ω_b / 10^{-4}	Ω_{DE} / 10^{-3}	n_s / 10^{-4}	σ_8 / 10^{-3}
Planck + Stage I	5.8	1.5	6.1	36.9	7.0
Planck + Stage II	5.1	1.4	5.3	32.8	6.0
Planck + Facility	2.7	1.2	2.7	21.9	3.3
Planck + DETF IV	2.2	1.0	2.2	16.0	2.3
Planck + WMAP	12	2.8	17	73.0	12
Planck+WP+BAO	7.8	2.5	10	57.0	11

TABLE 3
FORECAST 1σ MARGINAL ERRORS ON VANILLA Λ CDM MODEL PARAMETERS FOR THE SET OF REFERENCE SURVEYS, COMPARED WITH CURRENT CONSTRAINTS FROM PLANCK (TEMPERATURE-ONLY) AND WMAP (PLANCK COLLABORATION 2013B).

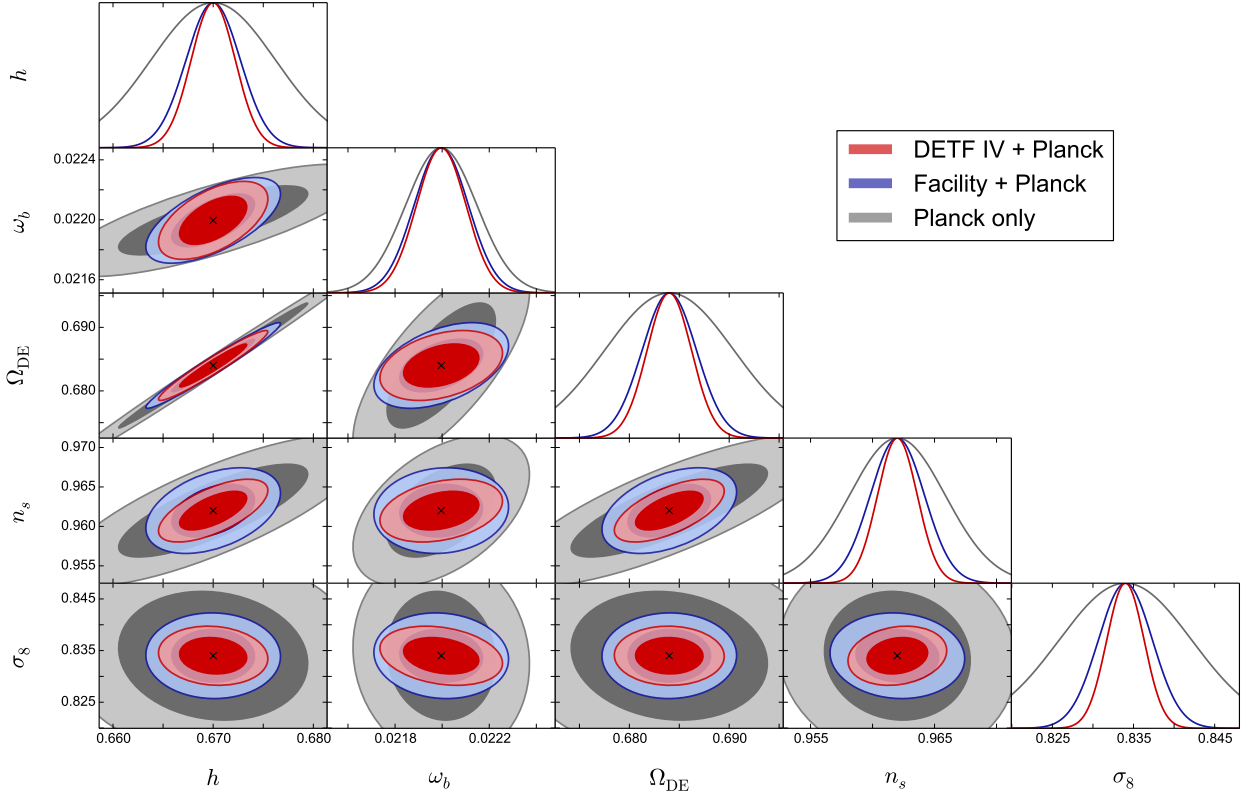


FIG. 11.— Forecasts for a six-parameter Λ CDM model (the sixth parameter, τ , was marginalised over in advance in the Planck Fisher matrix). This model has fixed $\Omega_K = 0$, $w_0 = -1$, and $w_a = 0$.

with IM (and redshift surveys). This is contingent on the assumption of a fixed $w = -1$, however.

The biggest effect of adding IM data (or indeed any LSS data) is to substantially improve the constraints on h and Ω_{DE} . As shown in Fig. 11, the two are strongly correlated for Planck alone, as the CMB only measures the combination $\Omega_M h^3 = (1 - \Omega_{\text{DE}})h^3$ (Planck Collaboration 2013b). The distance measures at late times depend on different combinations of these parameters and so help to break this degeneracy. This has a knock-on effect on other parameters that are correlated with them, especially n_s .

Table 3 summarises the vanilla Λ CDM constraints for the full set of reference experiments. The majority of future HI surveys are capable of improving on constraints from Planck plus existing LSS datasets; notably, H_0 and Ω_{DE} are determined at the sub-1% level by all but the smallest IM experiments.

High-end Facility-class experiments should even be competitive with a DETF Stage IV galaxy survey (c.f. Euclid or LSST), although this is only likely to be the case for parameters constrained by the distance measures, such as Ω_{DE} . Those that depend more on the power spectrum at smaller scales will not be quite as close, because the galaxy surveys measure $P(k)$ significantly better at relatively high wavenumbers of $k \sim 0.1 \text{ Mpc}^{-1}$ (Fig. 4); our IM experiments fall behind on these scales due to limited (single-dish) angular resolution.

5.2. Dark energy equation of state

The driver for the majority of the cosmological surveys currently under development is to find precision constraints on the dark energy equation of state, $w(z)$, and in doing so to infer the physical nature of the substance that appears to be driving the accelerated expansion of the Universe. HI intensity mapping provides a way of constraining $w(z)$ with considerable precision, using the full combination of $D_A(z)$, $H(z)$, and $f(z)$ reconstructed over a broad range of redshifts.

While the evolution of the equation of state parameter depends on the underlying dark energy theory, and as such could take any number of functional forms, it is nevertheless useful to work with a simple expansion about $z = 0$,

$$w(a) \approx w_0 + \frac{z}{1+z} w_a. \quad (21)$$

This commonly-used parametrisation should be reasonably accurate at late times, but will not capture more exotic behaviour at $z \gg 1$. The corresponding dark energy density evolves with redshift as

$$\Omega_{\text{DE}}(z) = \Omega_{\text{DE},0} \exp[3w_a z/(1+z)] (1+z)^{3(1+w_0+w_a)}. \quad (22)$$

The overall sensitivity of an experiment to a varying equation of state can be summarised (to some extent) by the dark energy *figure of merit*, defined by the Dark




FIG. 12.— Effect of various priors on $w_0 - w_a$ constraints, for Facility + Planck. Ω_K is already well-constrained by the combination of CMB and HI data, so the flatness prior has only a small effect. Additional H_0 information has a larger effect in breaking the degeneracy.

Energy Task Force as (Albrecht et al. 2009)

$$\text{FOM} = 1/\sqrt{\det(F^{-1}|_{w_0, w_a})}, \quad (23)$$

which is proportional to the reciprocal of the area enclosed by the 68% contour of the (w_0, w_a) error ellipse for Fisher matrix F .

The foremost task in understanding the nature of dark energy is to determine whether the equation of state differs from that of a cosmological constant, $w = -1$. Current constraints on w_0 and w_a are relatively weak; the combination of Planck with SNLS supernova data does give values that are slightly in tension with a pure cosmological constant (Planck Collaboration 2013b), but the significance fades when other datasets are used instead. Fig. 13 shows the improved constraints that can be expected on w_0 , w_a , and Ω_{DE} for the combination of our reference experiments with Planck, assuming flatness. Despite the addition of IM data, the parameters remain strongly correlated, so even substantial deviations from $w = -1$ will not necessarily be picked up. Nevertheless, a substantial fraction of the $w_0 - w_a$ plane can be excluded by IM + Planck, so a successful detection is still possible if the real values lie orthogonal to the degeneracy direction. 1D marginal constraints for the full set of extensions to Λ CDM that we are considering here (including w_0 and w_a) are given in Table 4 for all of the experiments from Sect. 2.5.

If one takes the possibility of a varying equation of state seriously, w_0 and w_a should be left free when deriving constraints on other cosmological parameters. Table 4 shows the effect of marginalising over the equation of state on the vanilla Λ CDM model parameters. The parameters derived from the various distance measures are strongly affected – their 1D marginal uncertainty is typically increased by around an order of magnitude compared to the unmarginalised case shown in Table 3. This can be understood in terms of the degeneracies shown in Fig. 13; adding new parameters always increases the overall uncertainty, but because Ω_{DE} (and h) are highly correlated with w_0 and w_a , they are particularly strongly affected. Parameters that do not depend on distance




FIG. 13.— Top Panel: Forecast constraints on w_0 and w_a , including the Planck prior. We have assumed flatness ($\Omega_K = 0$), and fixed γ to its fiducial value. The DETF figures of merit for the Stage II, Facility, and DETF Stage IV surveys are 95, 358 and 712 respectively. Bottom Panel: Forecast constraints on w_0 and Ω_{DE} for the same setup.

measures, i.e. n_s and σ_8 , are less affected by the equation of state parameters, and so their marginal uncertainties increase by only a modest amount.

As we have seen, even the addition of intensity mapping or other intermediate-redshift LSS data to the CMB constraints is insufficient to break all of the parameter degeneracies once w_0 and w_a are allowed to vary. In order to precisely determine these parameters, it is therefore necessary to add more data. Distance measurements from Type Ia supernovae are the obvious candidate, since they offer orthogonal constraints on $\Omega_{\text{DE}} - \Omega_M$ (Efstathiou & Bond 1999). A local measurement of H_0 is also useful; as shown in Fig. 11, h is strongly correlated with the dark energy density, so additional information about either parameter can substantially improve the constraints on both. Fig. 12 shows the effect of adding H_0 data to Planck + Facility. We also consider the effect of allowing departures from spatial flatness; as we will see in the next section, the combination of CMB and intensity mapping data measure Ω_K well, mostly independent of dark energy, so marginalising over curvature has a relatively minor effect on the $w_0 - w_a$ ellipse.

Fig. 14 shows the contribution to the dark energy figure of merit from each redshift bin. For our reference IM experiments, it is clear that the redshift range $z \lesssim 1.2$ is most critical; little improvement in FOM is seen above this redshift. The same cannot be said for the galaxy survey, however, which sees a roughly equal increase in FOM with each additional redshift bin across its whole z range. One way of understanding this behaviour is to compare Fig. 14 with the plots for $D_A(z)$, $H(z)$, and $f\sigma_8(z)$ in Fig. 6. Above $z \sim 1.2$, the angular diameter

Experiments	A / 10^{-2}	h / 10^{-3}	Ω_K / 10^{-4}	Ω_{DE} / 10^{-3}	n_s / 10^{-4}	σ_8 / 10^{-3}	γ / 10^{-2}	w_0 / 10^{-2}	w_a / 10^{-2}	FOM
Stage I	18.9	32.3	47.9	22.3	38.5	8.1	3.6	31.3	85.2	13.8
Stage II	13.2	23.7	33.1	17.2	38.0	7.8	4.4	15.2	33.1	39.9
Facility	5.2	8.7	13.6	6.9	35.0	6.0	1.8	5.4	14.9	265.4
GBT	73.9	131.9	178.4	93.4	38.6	8.2	20.1	95.0	221.8	1.1
GBT-HIM	31.2	64.3	78.9	45.4	38.6	8.2	9.8	50.7	126.9	4.2
GMRT	54.3	37.1	153.0	19.3	38.5	8.2	4.1	35.8	184.2	7.0
JVLA	57.7	43.0	175.3	22.6	38.6	8.2	4.5	40.1	209.2	5.5
Parkes	51.2	28.4	322.6	32.6	38.4	8.2	2.7	44.7	335.5	3.6
VLBA	74.8	47.8	826.9	86.2	38.6	8.2	3.8	91.7	799.8	0.8
WSRT + APERTIF	11.2	11.1	41.2	6.5	37.7	8.0	1.5	15.1	66.4	57.6
BAOBAB-128	24.3	50.2	71.3	36.6	38.5	8.1	9.0	33.3	71.4	8.0
BINGO	25.8	30.8	90.0	16.1	38.5	8.2	2.8	44.1	172.5	7.8
CHIME	3.0	8.7	9.7	7.1	30.2	5.2	3.4	5.0	15.1	288.1
FAST	7.5	13.5	16.0	10.1	33.5	6.4	3.2	7.1	18.5	144.7
MFAA	5.7	11.9	14.1	9.1	32.2	6.0	3.1	6.3	17.2	165.7
Tianlai	3.6	8.0	11.9	6.3	28.7	4.9	2.4	4.0	12.0	383.3
ASKAP	7.7	16.2	21.1	11.9	37.8	7.7	2.9	11.8	26.8	80.3
KAT7	114.0	76.4	1182.5	124.1	38.6	8.2	5.8	130.1	1138.6	0.4
MeerKAT (B1)	12.2	24.4	29.4	17.9	38.1	7.9	3.6	17.4	38.4	35.9
MeerKAT (B2)	10.2	9.4	26.8	6.1	37.5	7.7	1.5	6.4	29.5	171.4
SKA1-MID (B1) Autocorr.	6.2	11.2	16.1	8.7	35.9	6.6	2.3	7.1	17.6	162.5
SKA1-MID (B1) Interferom.	22.3	29.1	34.3	19.9	37.2	7.8	8.7	13.6	33.8	45.1
SKA1-MID (B2) Autocorr.	7.6	7.1	18.6	5.1	35.9	7.2	1.3	3.6	16.4	410.9
SKA1-MID (B2) Interferom.	368.2	37.3	94.3	19.0	38.0	8.2	10.6	22.8	86.5	18.5
SKA1-SUR (B1)	5.3	11.9	15.2	9.3	35.4	6.7	3.3	6.5	16.0	159.5
SKA1-SUR (B2)	4.5	6.5	12.2	5.3	35.3	5.7	1.4	3.8	12.2	444.2
SKA1-MID + MeerKAT (B1)	6.4	11.6	16.7	9.0	36.1	6.8	2.4	7.5	18.2	148.9
SKA1-MID + MeerKAT (B2)	7.7	7.1	18.6	5.1	35.9	7.2	1.3	3.5	16.3	414.7
DETF Stage IV (gal. survey)	2.4	7.5	8.6	6.2	27.1	5.3	3.2	4.1	12.8	405.5
Fiducial values	1.0	0.67	0.0	0.684	0.962	0.834	0.55	-1.0	0.0	—

TABLE 4

1D MARGINAL CONSTRAINTS (68% CL) ON THE EXTENDED Λ CDM MODEL, INCLUDING THE PLANCK PRIOR. THE CONSTRAINT ON A (WHICH HAS BEEN SUMMED OVER ALL REDSHIFT BINS) GIVES A MEASURE OF THE DETECTABILITY OF THE BAO.




FIG. 14.— Improvement in dark energy FOM as a function of the maximum redshift of the survey (Ω_K and γ marginalised).

distance and growth rate constraints begin to worsen for Facility, but remain relatively flat for the galaxy survey. Since w_0 and w_a are obtained from projections of these functions, it is no surprise that little is gained on the

FOM at redshifts where they are poorly constrained.

5.3. Curvature

The potential for HI intensity mapping experiments to span extremely wide redshift ranges – from $z \approx 0.1$ out to $z \gtrsim 2.5$ without too much difficulty – makes them an interesting prospect for unravelling the geometric degeneracy, i.e. the interplay between dark energy and curvature. Without strong assumptions on one or the other, it is difficult to separate the effects of Ω_K and $w(z)$ using only a single type of distance measure (Mortonson 2009; Shafieloo & Linder 2011), and for the CMB power spectrum alone they are completely degenerate. As discussed in Sect. 4.2, intensity mapping provides a suite of distance measures. The combination of IM and (e.g.) CMB data should therefore be very useful in separating curvature from the evolution of dark energy in a precise and unambiguous manner.

A precision determination of spatial curvature on horizon scales would also provide a rare opportunity to test inflation. Current observations seem to point in the direction of flatness, with the most recent bounds from Planck finding $|\Omega_K| \lesssim 10^{-2}$ (95% CL), consistent with the vast majority of inflation models, but if a detection of $|\Omega_K| \gtrsim 10^{-4}$ were made, the whole class of eternally




FIG. 15.— Improvement in Ω_K constraints as a function of maximum redshift of the survey. We have marginalised over w_0 , w_a and γ here.

inflating models would be put under pressure (Kleban & Schillo 2012; Guth & Nomura 2012).

The minimum curvature that can be detected unambiguously also happens to be at around the 10^{-4} level (Vardanyan et al. 2009; Bull & Kamionkowski 2013). Future CMB experiments should be able to approach this order of magnitude if w is fixed to -1 , and so too should a Facility-type IM experiment combined with Planck, as shown in Fig. 16. There is little justification for putting such a strong prior on the equation of state, though – any rigorous constraint on Ω_K must confront the geometric degeneracy head-on, and allow the full freedom of $w(z)$. Fig. 16 also shows the limits on curvature that can be achieved when the equation of state is left free. Though clearly worse than for fixed $w = -1$, the difference is relatively modest – the combination of a Facility survey with Planck should still be able to measure $|\Omega_K|$ to around 10^{-3} without any particularly strong assumptions about the form of $w(z)$.

The effect of the geometric degeneracy runs both ways – a lack of knowledge about Ω_K also degrades the reconstruction of the time evolution of the equation of state. Indeed, a percent level uncertainty in Ω_K can lead to a $\sim 100\%$ error on the recovery of more exotic forms for $w(z)$ (Clarkson et al. 2007), although it has been argued that current constraints can already mitigate this (Okouma et al. 2013). As shown in Fig. 12, the errors on the equation of state parameters do increase when we allow Ω_K to be free, albeit not substantially in the case of both Stage II and Facility; the combination of $H(z)$, $D_A(z)$, and $f(z)$ measurements from IM, plus the Planck prior, is enough to prevent any strong degeneracies from completely killing the $w_0 - w_a$ constraints. In fact, Fig. 12 shows that they are more sensitive to assumptions about H_0 than to curvature.

It is improved knowledge of the late-time expansion that most helps separate the effects of curvature and dark energy. We see this clearly in Fig. 15, where $[\sigma(\Omega_K)]^{-1}$ is plotted as a function of the depth of each survey. There is an optimal point beyond which little new information is gained by the IM surveys, coinciding with the redshift at which the constraints on $f(z)$ and $D_A(z)$ start to degrade




FIG. 16.— Forecast marginal constraints on Ω_K (68% CL) for the reference experiments plus Planck, with (w_0, w_a) fixed to their fiducial values (upper errorbars), and marginalised over (lower errorbars). The shaded area shows the current best constraint on Ω_K (w_0, w_a fixed) from Planck + WMAP + high- ℓ CMB + BAO (Planck Collaboration 2013b).

due to the limited angular resolution of the experiments (Fig. 6). This happens at higher z_{\max} for the galaxy survey, which makes up for its lack of low redshift bins by having relatively flat fractional errors in $D_A(z)$, $H(z)$, and $f(z)$ out to $z \approx 2$. At even higher redshift, the dynamical effect of curvature is completely negligible, so little extra information could be gained anyway.

5.4. Parametrised growth history

The growth history of the Universe is a particularly powerful test of gravity. Modified theories of gravity generically alter the growth of structure from its GR behaviour, typically enhancing clustering on non-linear scales, and increasing peculiar velocities. Signatures of modified gravity in the non-linear regime are difficult to disentangle from less exotic astrophysical effects, leaving the linear velocity field as, in some sense, the ‘cleanest’ modified gravity observable from large scale structure.

There is great variety in the effects that different modifications to gravity have on the linear growth history; the space of theories is complex, and has proved difficult to parametrise in a simple way (Hu & Sawicki 2007; Baker et al. 2013; Battye & Pearson 2012). For the purposes of illustration, we will fall back on one of the simplest parametrisations of growth, using the growth index formulation of Peebles (1980): $f(z) = \Omega_M^\gamma(z)$. Deviations from GR are captured, in part at least, by the difference in γ from its Λ CDM+GR value, $\gamma_{\text{GR}} \approx 0.55$.

Two notes of caution are necessary: Firstly, many modified gravity theories do not have growth histories that are well-described by a constant γ . Allowing γ to be a function of redshift can help (Linder & Cahn 2007; Ishak & Dossett 2009), but even then there are many cases where the growth rate becomes scale-dependent, or is otherwise poorly described by this parametrisation. Secondly, dark energy models that require no modifications to GR can also modify the growth history, often in a way that is well-described by making the growth index a function of the equation of state parameter, $\gamma(w)$ (Linder 2005; Gong et al. 2011b). We neglect this possibility

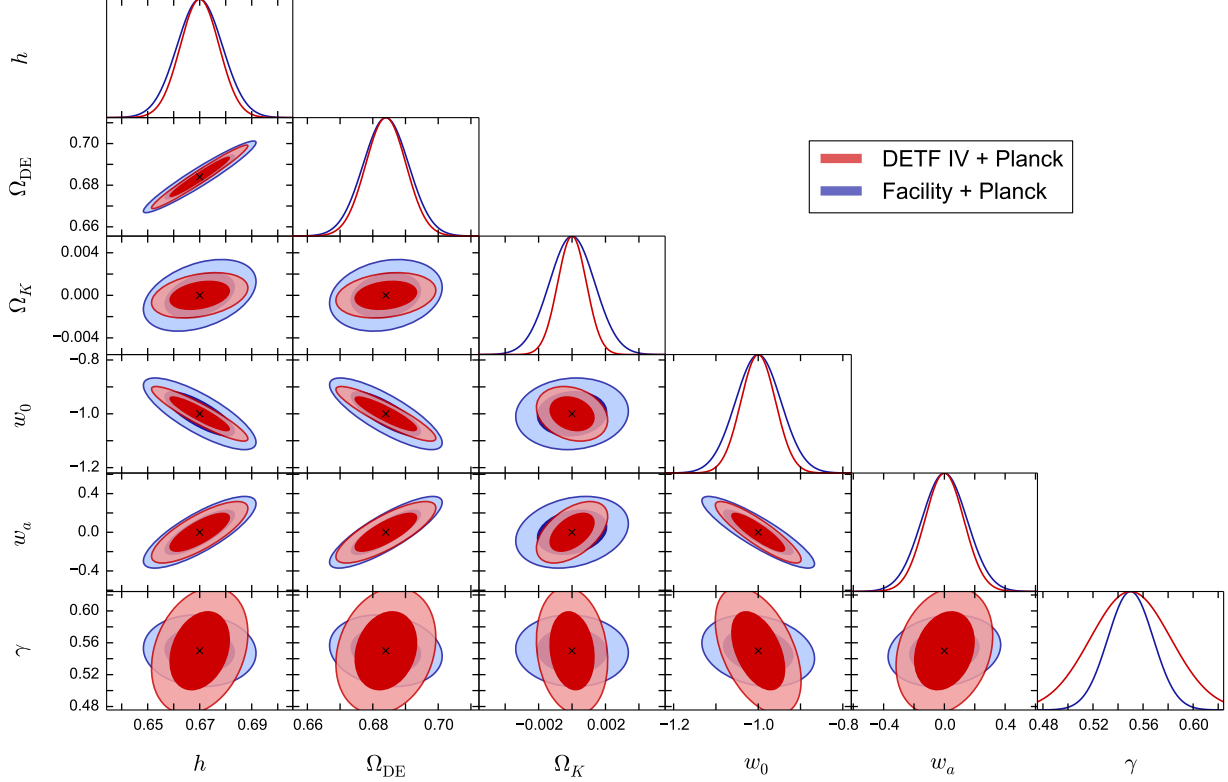


FIG. 17.— Forecasts for dark energy and modified growth parameters for two of the reference experiments. Note the significantly different behaviour with respect to the growth index, γ .

here, and instead treat γ as being independent of w .

Fig. 17 shows forecasts for the various dark energy parameters for the DETF Stage IV galaxy redshift survey and the Facility reference experiment. The 1D marginal constraints from the galaxy survey outperform the IM experiment for all parameters, except one – the growth index. Furthermore, the 2D constraints involving γ are roughly orthogonal between the two surveys, despite this not being the case for other combinations of parameters.

At first, this may seem surprising. In Fig. 6, the galaxy redshift survey constrains $f\sigma_8(z)$ to around 1% across most of its redshift range, while Facility’s precision can only match this in the lowest redshift bins, increasing to $\sim 4\%$ at higher z . For γ , though, it is the very lowest redshifts that make the most difference. At low z , the growth factor evolves most rapidly, and is most sensitive to the value of γ (i.e. $|df/d\gamma|$ increases as $z \rightarrow 0$), whereas at higher redshifts, matter begins to dominate, growth is slower, and the dependence on γ is relatively weak. By virtue of its substantially lower z_{\min} , then, the Facility experiment captures more of the redshift range most sensitive to γ , and wins out over the galaxy survey.

Fig. 18 shows the effect of the lowest redshift bins on γ more clearly. Even Stage II outperforms the Stage IV survey for $z_{\max} \lesssim 1$ – again thanks to its lower z_{\min} – despite producing significantly worse constraints on almost every other parameter. This behaviour is also related to the choice of distance measures, and how the degeneracies between them get broken. As shown in Fig. 8, the choice of measure can have a big effect on the $f(z)$ errors, so one might expect the strength of the constraint




FIG. 18.— Improvement in γ constraints a function of maximum redshift of the survey. We have marginalised over w_0 , w_a and Ω_K here. The low value in the first redshift bin for the IM experiments is due to a degeneracy.

on γ , and its orthogonality to the galaxy redshift survey, to change if a different subset of measures was used.

Assuming the full set of distance measures, the complementarity between intensity mapping and galaxy redshift surveys can be used to significantly increase the precision of the constraint on γ , to the point where it becomes possible to clearly distinguish between many modified gravity models. Fig. 19 shows the result of combining the two surveys on the errors for w_0 and γ , along with some




FIG. 19.— Constraints on (γ, w_0) for Facility, the DETF Stage IV survey, and the combination of the two, including Planck CMB priors. w_a and Ω_K have been marginalised over, and the biases for both surveys, $b_{\text{HI}}(z)$ and $b_{\text{gal}}(z)$, are free in each bin. Also plotted are example $f(R)$ and DGP modified gravity models from Amendola et al. (2013).

example predictions from modified gravity theories. The marginal error on γ goes from $\sigma_\gamma = 0.024$ for Facility + Planck to $\sigma_\gamma = 0.015$ for Facility + DETF IV + Planck.

6. SYSTEMATIC EFFECTS

Throughout this paper, we have compared the results from IM mapping experiments with those of a DETF Stage IV spectroscopic survey. In fact we have gone further and argued, in Section 3, that we can think of an IM survey as a spectroscopic survey with an anisotropic $V_{\text{eff}}(k)$, with very distinctive characteristics tied to the chosen mode of operation (single-dish or interferometric). While a useful comparison, IM experiments have their own, particular types of systematics that must be dealt with. In what follows we touch on these effects and attempt to quantify their impact on the target science.

6.1. Evolution of the cosmological HI signal

We have assumed a fiducial value of $\Omega_{\text{HI},0} = 6.5 \times 10^{-4}$ throughout our analysis. Clearly our results will strongly depend on this value, as it is important in setting the “signal-to-noise” of the experiment – the more neutral hydrogen there is, the more easily the cosmological signal can be detected. There are, however, large uncertainties in $\Omega_{\text{HI}}(z)$ from current observations (Fig. 20). In particular, different tracers of the HI density give inconsistent results, so neither the normalisation nor the redshift evolution of $\Omega_{\text{HI}}(z)$ are well understood.

The constraint of most relevance to us is from Masui et al. (2013), where IM measurements were cross-correlated with the WiggleZ galaxy redshift survey. It was found that

$$\Omega_{\text{HI}} b_{\text{HI}} \bar{r} = 4.3 \pm 0.7 (\text{stat.}) \pm 0.4 (\text{sys.}) \times 10^{-4}$$

at $z = 0.8$, where the best theoretical estimates for the cross-correlation coefficient are $\bar{r} = 0.9 - 0.95$. This was obtained by restricting the analysis to $0.075h \text{ Mpc}^{-1} < k < 0.3h \text{ Mpc}^{-1}$; extending it to $0.04h \text{ Mpc}^{-1} < k <$




FIG. 20.— Current constraints on the HI density fraction as a function of redshift, $\Omega_{\text{HI}}(z)$ (Meiring et al. 2011; Prochaska & Wolfe 2009; Noterdaeme et al. 2009; Rao et al. 2006; Martin et al. 2010; Lah et al. 2007; Zwaan et al. 2005; Khandai et al. 2011), partially based on the compilation in Duffy et al. (2012) (see also Padmanabhan et al. (2014)). DLA observations are shown in blue, cross-correlations in yellow, other observations in red, and simulations in green. The thick black line shows the fiducial $\Omega_{\text{HI}}(z)$ that we have adopted in this paper, which has $\Omega_{\text{HI}}(z=0) = 4.86 \times 10^{-4}$.

$0.8h \text{ Mpc}^{-1}$ lowers the constraint slightly to

$$\Omega_{\text{HI}} b_{\text{HI}} \bar{r} = 4.0 \pm 0.5 (\text{stat.}) \pm 0.4 (\text{sys.}) \times 10^{-4}.$$

The value of Ω_{HI} is entangled with the bias which, from semi-analytic models combined with N-body simulations (Khandai et al. 2011), is found to be consistently low, $b_{\text{HI}} \approx 0.55 - 0.66$, although it can go up to unity for certain model choices. As a result, (Masui et al. 2013) proposes that one should assume $\Omega_{\text{HI}} = 4.5 - 7.5 \times 10^{-4}$ at $z = 0.8$.

The lack of agreement between observations makes it difficult to reconstruct the redshift evolution of Ω_{HI} . At the upper end of the redshift range we are considering ($z \lesssim 3$), constraints from Damped Lyman- α (DLA) systems are scattered between $\Omega_{\text{HI}} \approx 4 - 9 \times 10^{-4}$. At $z \sim 1$ there are discrepant results between theoretical models that find $\Omega_{\text{HI}} \approx 3 \times 10^{-4}$ (Duffy et al. 2012) and observations of DLAs with HST that give $\Omega_{\text{HI}} \approx 9 \times 10^{-4}$. At $z = 0$, the ALFAFA and HIPASS surveys find $\Omega_{\text{HI}} \approx 4 \times 10^{-4}$, which is slightly lower than our fiducial $\Omega_{\text{HI},0}$.

For the forecasts in this paper, we tread the middle ground (solid line, Fig. 20). The $\Omega_{\text{HI}}(z)$ redshift evolution is derived from a simulated HI halo mass function, as described in App. B, and we choose its fiducial normalisation to be consistent with the GBT/WiggleZ cross-correlation measurement at $z = 0.8$. The magnitude and redshift evolution of the HI bias, $b_{\text{HI}}(z)$, are derived from the same mass function.

Fig. 21 shows the effect of rescaling Ω_{HI} by a constant factor on the constraints for several observables. If $\Omega_{\text{HI}}(z)$ was halved, for example, the FOM for Facility would drop by a factor of 5. This highlights the sensitivity of cosmological constraints from IM to the HI density, and gives some idea of the degradation/improvement in




FIG. 21.— Normalised FOM and Ω_K and γ marginal errors, as a function of Ω_{HI} rescaled by a constant factor, for the Facility (blue, solid) and Stage II (green, dashed) surveys.

performance that would be expected if $\Omega_{\text{HI}}(z)$ substantially differs from what we have assumed.

6.2. Non-linear scale

We have marginalised over the non-linear scale, σ_{NL} , in all of our forecasts. As described in Section 3, this parameter is responsible for setting the resolution in the radial direction; beyond this scale, non-linear peculiar velocities wash out all redshift information.

Fig. 22 shows the effect of changing the fiducial non-linear scale. As one might expect, increasing σ_{NL} degrades the various constraints, as information is lost at progressively larger scales. For the Facility experiment, which uses a combined single-dish and interferometer mode, the change in non-linear scale has a similar, relatively mild effect on all of the figures of merit, which change by less than a factor of two for a doubling of σ_{NL} . This is not the case for the purely interferometric Stage II survey, which is more sensitive to smaller scales (especially at low z), so is hit harder by the loss of information there.

6.3. Foreground contamination

The viability of intensity mapping as a cosmological probe vitally depends on the availability of accurate foreground removal techniques, as the contaminating signals have an amplitude of between 4 to 6 orders of magnitude greater than the cosmological HI signal (e.g. Alonso et al. (2014)). In the previous sections we adopted a fiducial value for the residual foreground contamination amplitude of $\epsilon_{\text{FG}} = 10^{-6}$, which is a reasonable target value for current foreground subtraction methods. In this section we quantify the sensitivity of our forecasts to the assumed removal efficiency, and discuss a number of potential problems surrounding foreground contamination.

Most of the proposals for how to subtract foregrounds from IM data rely on a simple qualitative assumption: that foregrounds have a smooth (coherent) frequency dependence over the observed frequency ranges⁵. This is generally true, apart from in the presence of polarisation

⁵See Morales et al. (2006) for foregrounds affecting the Epoch of Reionisation.




FIG. 22.— Normalised FOM/marginal errors as a function of σ_{NL} (see Fig. 21 for key).

leakage, which we will discuss shortly. A simple common-sense approach is to assume that the foreground signal along the frequency direction is accurately modelled as a sum of low-order polynomials or similar, which capture what should essentially be a (very mildly) modulated power law behaviour. This is at the heart of the methods presented in Wang et al. (2006); Gleser et al. (2008); Jelić et al. (2008); Liu et al. (2009).

Another possibility is to be agnostic about the frequency dependence of the foregrounds, but decompose the total signal in some form of signal-to-noise eigenbasis. Since the foregrounds have such large magnitudes, the hope is that they will be contained only in the very high signal-to-noise part, and thus will be suitably segregated from the cosmic signal. This is the logic behind the methods used in Chang et al. (2010); Wolz et al. (2013). When applied to real data in Chang et al. (2010), it was found that the foregrounds were not as strongly segregated from the cosmic signal as expected, so it was necessary to subtract a larger number of modes than originally planned (inevitably throwing-out some of the cosmological signal too). Even in the optimal case, this type of foreground removal method has been shown to leave a residual bias in (e.g.) the cosmological parameters that best-fit the recovered BAO (Wolz et al. 2013).

As described in Section 2.3, we model the effects of foregrounds in terms of a residual noise term with an overall relative amplitude, ϵ_{FG} , and a minimum cutoff wavenumber, k_{FG} , along the radial (frequency) direction. For any given foreground removal method, these two parameters are intertwined – the more large-scale modes one uses to estimate the shape of the foregrounds (i.e. the larger k_{FG} is made), the better the removal efficiency, and so the lower ϵ_{FG} should be. We have chosen k_{FG} to be a fixed fraction of the total bandwidth across *all* the redshift slices of a survey, and have implicitly absorbed the removal efficiency into ϵ_{FG} . (Note that this model is purely stochastic, and does not allow us to model the biases that were discussed in Wolz et al. (2013).)

In Fig. 23, one can clearly see that the optimal level of foreground subtraction is around $\epsilon_{\text{FG}} \simeq 10^{-6}$ for the Facility configuration, but can be larger for other configurations (which have higher noise levels). The impact of




FIG. 23.— Normalised FOM/marginal errors as a function of ϵ_{FG} (see Fig. 21 for key).

changing k_{FG} is shown in Fig. 24; while larger k_{FG} should yield better foreground subtraction, there is a trade-off involved in losing more data at small k .

While our analysis has so far focused on unpolarised foregrounds, problems may arise if one considers instrumental leakage from *polarised* foregrounds into the total intensity mode. For a typical receiver, one expects a cross-leakage of order a few percent, so given the large amplitude of the foregrounds, this can have a significant effect on the total signal. Synchrotron emission is the main polarised foreground, and has a non-trivial angular and frequency dependence due to Faraday rotation. To see this, consider the polarisation angle, ϕ , which is rotated by the galactic magnetic field, \mathbf{B} , through

$$\phi(r) = \phi_0(r) + c^2 \nu^{-2} \int_0^r n_e(r') \mathbf{B}(r') \cdot d\mathbf{r}', \quad (24)$$

where $\phi_0(r)$ is the initial angle, n_e is the electron density, and r is the distance along the line of sight. The galactic magnetic field is a non-linear superposition of an overall coherent mode, tied to the spiral structure of the Galaxy, with a turbulent stochastic mode on small scales (Beck 2001). From (24), one can see that the rotation of the polarisation vector depends on both frequency and the line-of-sight through the galaxy with some decoherence length (corresponding to the coherence of \mathbf{B} along the line of sight). This leads to a more complex foreground signal that is considerably less smooth in frequency than unpolarised foregrounds – an effect that increases in severity at low frequencies (Alonso et al. 2014). While we have not explicitly included it in our foreground model, the effects of polarisation leakage can be partially accounted for by a larger fiducial ϵ_{FG} .

6.4. Autocorrelation calibration (single-dish)

In this paper, we have advocated using some instruments as collections of single-dish experiments, i.e. in autocorrelation mode. This is common practice in CMB mapping experiments, and has been the leading method of producing large-scale, high-resolution maps. Its use with radio telescope arrays at lower frequencies is less common, however, and must be treated with some care because of a number of potentially serious systematics.




FIG. 24.— The effect of rescaling the foreground cutoff scale, k_{FG} , on the normalised FOM/marginal errors (see Fig. 21 for key). The base value of the cut-off scale is $k_{\text{FG},0} = 2\pi/(r_\nu \Delta \tilde{\nu}_{\text{tot}})$.

Nevertheless, there is some precedent for using autocorrelation mode to detect both individual HI sources and unresolved emission. In Braun et al. (2003), the WSRT array was used in this mode to perform a wide-field survey, yielding a sample of ~ 150 HI galaxies. A key difficulty of the analysis was in obtaining an accurate calibration of the autocorrelation mode – while the average gain over all receivers was relatively stable, there were variations of up to 10% for individual receivers. This was calibrated out by using cross-correlation data for known radio sources. For the first attempt at mapping the unresolved HI signal with the GBT telescope (Chang et al. 2010), the flux calibration was controlled by fixing an intermittent noise source at the feed point, and by periodically monitoring a known source.

Drifts in the gain (e.g. due to instrumental temperature variations) are just one type of autocorrelation systematic. Another is due to spillover and sidelobe pickup, which can arise from a poorly characterised beam and ground contamination. This is not an insurmountable problem, and ground-based and balloon-borne CMB experiments commonly incorporate design features to mitigate these effects. One approach for HI instruments is that proposed by BINGO (Battye et al. 2013). There, the idea is to use a partially-illuminated aperture to reduce the effect of sidelobes, spillover, and RFI contamination. Another aspect of the BINGO design is the use of a fixed dish, which scans the sky simply by allowing it to drift through the beam as the Earth rotates. **This bypasses various problems that arise with moving parts, and allows a more precise pointing calibration than is possible with dynamic “raster” scan strategies.**⁶

Another key obstacle in the analysis of autocorrelation data is the presence of $1/f$ noise, which is a coherent (correlated) noise drift on long timescales. Ideally, one would be able to construct a receiver system such that the $1/f$ knee (i.e. the timescale beyond which correlations become important) is at very low frequencies – a few $\times 10^{-3}$ Hz, for example. If this is possible, then the noise will drift over periods of several minutes which,

⁶See also the ‘on-the-fly’ technique (Mangum et al. 2007).

for a drift scan, corresponds to angular scales of a few degrees (i.e. larger than the BAO scale). This is the method used by BINGO.

More traditionally, the way to deal with noise drifts or biases in CMB experiments has been to devise a scan strategy such that the modes one is looking at are scanned at frequencies higher than the knee frequency. The signal in a given pixel will then be localised in particular (frequency) modes of the time series which are subject only to white (i.e. uncorrelated) noise. This requires that the instrument should have the ability to scan quickly across the sky, which can be challenging for large dishes. One can further mitigate the effect of $1/f$ noise by devising a scan strategy with multiple cross-linking, i.e. in which each pixel is revisited a number of times on different time scales, from different directions. The inversion process for extracting the map from the time series (and estimating the noise) is then well-defined and numerically robust (Ferreira & Jaffe 2000). If the noise is smooth in frequency, it may also be possible to clean out any noise bias component, as one would a foreground component in the signal.

In summary, there are clearly a number of issues that must be considered carefully when working in autocorrelation mode, but as we have shown in our analysis throughout this paper, the scientific potential of single-dish IM experiments is tremendous. Ideally, one would design the dishes, receivers, and other hardware of a given array specifically to mitigate the problems discussed above, but in some cases (e.g. with the SKA) this is not possible due to competing design constraints enforced by the need to operate primarily in an interferometric mode – optimising the hardware for autocorrelation mode as well is simply too expensive. Any chance of controlling these (potentially critical) systematic effects to a sufficient degree is then left down to the choice survey strategy and data analysis methodology.

Experience with autocorrelation CMB experiments suggests that reliance on non-hardware techniques to reduce important systematics is a risky strategy, but there are a number of features of forthcoming radio telescope arrays that may be helpful in making this feasible. For example, an experiment with hundreds of individual dishes can make use of the fact that some systematics will be uncorrelated between the dishes; by cross-correlating (detected) maps of the same volume produced by different dishes, many effects can therefore be expected to correlate out. Nevertheless, it remains to be demonstrated that an experiment as complex as Phase 1 of the SKA can successfully control its autocorrelation calibration down to the required level, and how the survey/analysis strategy affects its overall sensitivity. This must therefore be seen as an important caveat of our analysis.

6.5. Sensitivity to large scales (interferometers)

A particular limitation to HI mapping with interferometers is the difficulty of sampling modes on large angular scales. In the simplest model of an interferometer – as a collection of dishes – the minimum measurable wavenumber is set by the minimum baseline, which cannot be smaller than the diameter of the dishes. From Fig. 3 we can see that arrays with large dishes will not adequately sample BAO scales at low redshift in interferometer mode, so most of the constraints must come

from single-dish mode.

There are a few ways to mitigate this shortcoming. The simplest is the approach effectively taken by BAOBAB and dense aperture arrays – to just use smaller dishes and pack them closer together. This results in smaller baselines and a larger field of view for the interferometer, but reduces its total effective collecting area (and thus its sensitivity). One can add more dishes to compensate, although this can substantially increase the cost of correlation hardware, which scales roughly like $\sim N_{\text{dish}}^2$. GPU-based correlators, or correlating only a subset of receiver pairs, can reduce costs for large numbers of receivers.

Alternatively, one can use a more novel reflector design. For example, CHIME uses long cylindrical reflectors with many closely-spaced receivers installed along the cylinder (Shaw et al. 2014). This provides a large number of short baselines, and a primary beam that is $\sim 180^\circ$ in one direction but much narrower along the orthogonal direction. This leads to a very anisotropic sampling of transverse Fourier modes, and only the large angular modes that are exactly aligned with the cylinders will be properly sampled. Because the visibilities measured by the interferometer are convolved with a window function defined by the primary beam, however, modes larger than the shortest side of the FOV will be aliased, making them difficult to disentangle. This is only the case if the interferometer tracks a single patch of the sky, though; if one progressively scans over the patch with different pointing offsets, and has precise knowledge of the primary beam pattern, it is possible to remove the aliasing effect of the primary beam and thus independently measure modes larger than the instantaneous FOV by *mosaicing* (Holdaway 1999). In the case of CHIME, drift scanning provides a continuous range of pointing offsets, and in principle the array can see the whole sky over a 24 hour observation period.

One can also make interferometric measurements over a number of separate pointings without mosaicing, simply to survey a larger area of sky (White et al. 1999). Drift scanning can be seen as a continuous limit of this. The advantage of such a method is that one can greatly reduce the sample variance of the smallest-baseline modes, simply by observing them on several independent patches of the sky. A crucial point is that simply patching together multiple fields does not allow modes larger than those defined by the minimum baseline to be measured, and therefore does not change the range of modes sampled, but *does* increase the sensitivity within that range. We have implicitly assumed that interferometers can handle multiple pointings by allowing $S_{\text{area}} > \text{FOV}$ in our forecasts.

6.6. Combined mode

Another possibility is to operate some experiments in a “combined mode”, where both autocorrelation and cross-correlation data are collected. The simplest way of doing this in practise is to split the total survey time into two chunks, using only one of the observing modes for each. From the previous two sections, we can see that each mode will have different systematics and hardware requirements, and it is likely that substantially different survey strategies would be needed for each mode.

The situation is considerably more difficult if one tries




FIG. 25.— Normalised FOM/marginal errors as a function of survey duration, t_{tot} (see Fig. 21 for key).

to collect data in both modes simultaneously. As discussed above, in single-dish mode one has to mitigate $1/f$ noise, typically by rapidly scanning across the sky. Conversely, interferometers require precisely-measured baselines and pointings to allow accurate phase calibration and reconstruction of the beam pattern; as dishes accelerate while scanning, even small distortions of the mounts can make this difficult. The combination of all these issues can potentially be overcome by drift scanning or through the use of novel mounts, but neither option appears to have been tested yet.

7. AN IDEAL HI SURVEY?

In this section, we suggest what an ‘ideal’ future HI intensity mapping experiment for late-time cosmology would look like. So far, we have assumed that Phase I of the SKA will represent the pinnacle of HI intensity mapping science for the coming decade. While its performance in terms of cosmological parameter constraints will indeed be impressive, it is worth remembering that the SKA is a general purpose facility, and is not specifically designed for IM. We propose that a cheaper purpose-built instrument, optimised for HI science, would be able to match, and perhaps even surpass, the SKA’s performance.

For the sake of simplicity, we target the dark energy figure of merit (FOM) as the only parameter to be optimised for. Sensitivity to intermediate scales, where the BAO and other distance measures are most important, is therefore a priority, although distance information from RSDs and the overall shape of the matter power spectrum are also useful (see Sect. 4.2). Dark energy is typically most important at low redshift, and so one might reasonably expect to focus the survey on the interval $0 \leq z \leq 2$. The lowest redshifts are likely to be subject to RFI, however, and pushing to higher redshifts brings in issues of limited angular resolution and increasing galactic foreground emission.

In terms of the available technology, we assume that $T_{\text{inst}} \approx 25\text{K}$ wideband receivers can be built cheaply and in bulk, and that correlators for several thousand receivers will also be relatively affordable. Time allocation is not an issue for a purpose-built instrument and, as




FIG. 26.— Normalised FOM/marginal errors as a function of survey area, S_{area} (see Fig. 21 for key).

we can see from Fig. 25, one only gains from increasing the amount of integration time. We assume that at least 10,000 hours of effective observing time can be used. This leaves only a handful of basic design parameters that can be varied:

- Survey area, S_{area}
- Dish size, D_{dish}
- Array configuration (maximum and minimum baselines, D_{max} and D_{min} , and filling factor)
- Frequency range (corresponding to redshift range, $[z_{\text{min}}, z_{\text{max}}]$)

The optimal survey area for a fixed amount of integration time depends on how quickly one can integrate down to the signal-dominated regime at each pointing. In Fig. 26 we show that, in the case of the Stage II experiment, increasing S_{area} above its optimal value will lead to a reduction in overall signal-to-noise and hence in the overall FOM. This is not the case for the Facility experiment, which already has sufficient time to reach signal-domination at each pointing; increasing S_{area} simply reduces the cosmic variance and therefore improves the FOM. In designing an optimal survey, one should pick S_{area} such that it is signal-dominated at each pointing, but only just, so as not to spend too much time integrating in a regime that is dominated by cosmic variance. This is the approach used when designing survey strategies for CMB experiments.

For interferometers, a high filling factor is desirable, as it equates to higher sensitivity. For a fixed number of dishes, increasing the filling factor amounts to increasing the dish diameter, or decreasing the maximum baseline length. Smaller dishes are useful for increasing the field of view, and thus the survey speed, however, and allow for smaller minimum baselines, which is important for resolving larger angular scales, especially at low redshift.

For single-dish experiments, smaller dishes also improve survey speed due to their increased field of view, but this now comes at the cost of angular resolution. A balance must therefore be found between resolving intermediate scales over as much of the redshift range as

possible, and survey speed. There is also the issue that larger dishes cost more, but for a pure single-dish experiment this is offset by there being no need for expensive correlator hardware.

Under the design constraints that we imposed, it turns out that the SKA Phase I arrays are close to ideal for single-dish experiments, given the assumed foreground removal efficiency. This was to be expected; Facility (broadly representative of a single-dish SKA configuration) is already nearing the cosmic variance-limited constraints on $P(k)$ from the DETF Stage IV galaxy redshift survey, as we showed in Fig. 4. The only significant improvement to be had is around $k \sim 0.1 \text{ Mpc}^{-1}$, which could be obtained by increasing the survey time or, even more effectively, by decreasing z_{\min} . Shifting the maximum frequency of Facility from 1100 to 1200 MHz while keeping the total bandwidth fixed to 700 MHz effectively matches its FOM to that of the galaxy redshift survey.

As interferometers, the SKA configurations have too small a field of view and too low a filling factor to achieve competitive dark energy constraints. A purpose-built interferometer operating over the desired redshift range would be better off having much smaller dishes, closely packed together. This is the approach that CHIME and BAOBAB are effectively taking. A 250-element array with 2.5m dishes distributed over a 44m core (giving a filling factor of 0.8) would surpass Facility's FOM for $\nu_{\max} = 1100 \text{ MHz}$, and match the galaxy survey's for $\nu_{\max} = 1200 \text{ MHz}$ (where $\Delta\nu = 700 \text{ MHz}$ in both cases).

Is it possible for IM experiments to do better than the reference galaxy redshift survey? Yes, but not without relying on either a higher maximum frequency or small angular scales. On large scales, the single-dish SKA configurations are limited by residual foregrounds, as shown in Fig. 27. Extending to higher redshifts increases the total volume being probed, but the sensitivity to dark energy decreases significantly above $z \gtrsim 2$, so in practise little is gained by doing this. As we have already seen, going to lower redshift (i.e. increasing ν_{\max}) can have a big effect, as dark energy is most important here. The problem of RFI (radio interference) increases towards 1.4 GHz though, so this is also difficult. Besides, other sources of information on the matter density field are available at low- z (e.g. existing galaxy redshift surveys), so it is not clear whether extending IM surveys into this region would be particularly useful.

Another option is to improve sensitivity on small angular scales, $k_{\perp} \gg 0.1 \text{ Mpc}^{-1}$. This can be achieved by increasing the number of detectors, improving the single-dish angular resolution, reducing the instrumental noise, or performing longer surveys (Fig. 27 shows the ideal case). In theory this would provide extra distance information from the shape of the power spectrum, but this relies on being able to accurately model the non-linear power spectrum, which is also tricky. To significantly improve DE constraints past what a Facility-class experiment is capable of, one is probably better off focusing on combining IM with other probes, such as weak lensing.

8. DISCUSSION

Neutral hydrogen (HI) intensity mapping looks set to become a leading cosmology probe during this decade. In this paper we have assessed its potential for constraining cosmological parameters, focusing on 'late times', $z \lesssim 3$.




FIG. 27.— Fractional constraints on $P(k)$ for an ‘ideal’ (noise-free) 15m single-dish survey, covering the same redshift range and survey area as the DETF Stage IV reference experiment. The deviation from the cosmic variance limit on large scales is (partially) due to the k_{FG} cutoff.

We used a few reference experimental designs – Stage I, Stage II, and Facility – that are inspired by up-and-coming experiments to assess how well we will improve our knowledge of the standard cosmological model, the nature of dark energy, the spatial curvature of the Universe, and the growth rate of structure. We have done so being mindful of the potential systematic problems that need to be faced.

Intensity mapping at radio frequencies has a number of advantages over other large scale structure survey methodologies. Since we only care about the large-scale characteristics of the HI emission, there is no need to resolve and catalogue individual objects, which makes it much faster to survey large volumes. This also changes the characteristics of the data analysis problem; rather than looking at discrete objects, one is dealing with a continuous field, which opens up the possibility of using alternative analysis methods similar to those used (extremely successfully) for the CMB. Thanks to the narrow channel bandwidths of modern radio receivers, one automatically measures redshifts with high precision too, bypassing one of the most difficult aspects of performing a galaxy redshift survey.

These advantages, combined with the rapid development of suitable instruments over the coming decade, look set to turn HI intensity mapping into a highly competitive cosmological probe in only a short space of time. In Sect. 5, we showed that Facility-class experiments would be broadly competitive with DETF Stage IV galaxy redshift surveys such as Euclid, LSST, and WFIRST in terms of cosmological parameter constraints, in about the same timeframe. Indeed, the largest planned surveys, such as SKA1-SUR, may even be able to surpass the cutting-edge galaxy surveys, although this is contingent on the (currently poorly-known) HI density and the performance of foreground removal algorithms. Since the currently-planned Facility class surveys are not specifically designed for IM, we also considered what a large, purpose-built HI experiment would be able to achieve in Sect. 7. We found that little extra could be gained without pushing to higher frequencies or smaller




FIG. 28.— Constraints on w_0 and w_a from the combination of the DETF Stage IV galaxy survey and a combined-mode SKA1-MID configuration, compared with results for the experiments individually. The figures of merit are 427 (SKA1-MID), 438 (galaxy survey), and 2124 (combined). We have assumed that the survey volumes are independent; otherwise, cosmic variance would degrade the combined constraint.

(non-linear) angular scales; neither are free of problems.

More important than their individual performance is what IM and galaxy redshift surveys can do in combination. In Fig. 19, we showed that Facility and a DETF Stage IV survey give roughly orthogonal constraints on w_0 and γ when combined with CMB data, mostly as a result of their complementary redshift coverage. Fig. 28 shows the joint constraints on w_0 and w_a for the combination of DETF Stage IV and a combined-mode SKA configuration with a lower-redshift band; the resulting dark energy FOM is almost five times that of either survey individually. This large improvement is due to the increase in the total surveyed volume, as well as the complementary redshift coverage. One can also benefit from the “multi-tracer” effect, whereby the limits imposed by cosmic variance on some variables can be overcome by measuring several distinct populations of tracers of the cosmic density field (McDonald & Seljak 2009; Abramo & Leonard 2013). Combining HI intensity mapping and galaxy redshift surveys should therefore offer particularly stringent constraints on the dark energy equation of state and growth index parameters – an absolute necessity for distinguishing between different dark energy and modified gravity models.

HI intensity mapping experiments also offer some novel features – for example, in their ability to probe ultra large scales in the late Universe. Facility-class arrays like Phase I of the SKA will be able to simultaneously survey an extremely wide range of redshifts over greater than half of the sky, covering volumes of several tens of cubic Gpc in one fell swoop. This is sufficient to detect physical effects beyond the matter-radiation equality scale (Fig. 29), including non-Gaussianity, spatial curvature, and potential deviations from large-scale homogeneity and isotropy. As was shown in Fig. 3, a sufficiently large HI survey could even probe beyond the horizon size at $z \gtrsim 1$, allowing us to access causally-disconnected regions long after recombination.

Before HI intensity mapping can contribute seriously to




FIG. 29.— Forecast constraints on $P(k)$ from SKA1-MID (B1), assuming perfect foreground removal (i.e. $\epsilon_{\text{FG}} = k_{\text{FG}} = 0$). The turnover in the power spectrum should be clearly detectable.

late-time cosmology, a number of potential pitfalls must be navigated. Chief amongst these is the overall magnitude of the HI density, $\Omega_{\text{HI},0}$, and its evolution with redshift, both of which are currently poorly constrained. The lower the density, the harder the HI signal will be to detect (and the more aggressive the foreground cleaning will need to be). Fig. 21 showed the effect of changing $\Omega_{\text{HI},0}$ on various figures of merit; a factor of two reduction in HI density from our fiducial value results in roughly a factor of five degradation in parameter constraints, which would be troublesome, although not catastrophic. Similarly, a higher density would make the signal much easier to detect. This situation is loosely analogous to the dependence of galaxy cluster surveys on the normalisation of the power spectrum – when it was found that σ_8 was closer to 0.8 than 0.9, this vastly reduced expected cluster number counts, leading to a corresponding drop in forecast constraints from cluster surveys. For HI intensity mapping, all we can do is wait for better measurements of $\Omega_{\text{HI},0}$ to see what the effect will be.

A confounding factor that is more directly under our control is the foreground cleaning efficiency, which we investigated in Sect. 6.3. The galactic foreground signal is around six orders of magnitude larger than the cosmological HI signal, but has a distinctive (and thus easy to separate) behaviour for the most part. Since they vary on similar angular/frequency scales to the galactic foregrounds, large-scale cosmological modes are likely to be hit harder by imperfect foreground cleaning. There is no reason why this cannot be overcome, however – similarly large modes are routinely dealt with successfully in CMB analysis (Planck Collaboration 2013a). Of potentially more concern is the issue of polarisation leakage, which imprints a more variable signal on top of the cosmological one. Sufficiently sophisticated modelling, combined with a sustained effort to control or characterise leakage at the hardware level, should be able to deal with this.

Unsurprisingly, choosing the right survey strategy is vital; we investigated the effect of changing various survey parameters in Sect. 7. Throughout the paper, we have also considered the difference in performance between interferometer and single-dish observation modes. For purpose-built IM experiments (e.g. the ‘ideal’ ex-

periment described in Sect. 7), one will tend to prefer interferometry because of the comparative ease of controlling instrumental/atmospheric systematics, although this must be offset against the significant computational expense of correlating many baselines. For general-purpose instruments, whose design is likely to be set by other considerations, interferometry may be a poor choice – if the array has large dishes, its interferometric field of view will be small, making it relatively insensitive to the intermediate scales that are most useful for detecting the BAO (see Fig. 3). In this case, there is much to be gained by using a single-dish (or combined single-dish + interferometer) mode instead. This is the path that we advocate for Phase I of the SKA, at least for low redshifts – at higher redshifts, the angular resolution in single-dish mode is actually too low, so the errors on quantities such as $D_A(z)$ get larger (see Sect. 4.2). This choice brings with it a number of significant data analysis challenges however, as discussed in Section 6.4; how to precisely and consistently calibrate many hundreds of dishes operating in autocorrelation mode is currently an open problem, and the various advantages of single-dish operation will only be available if it can be solved.

As we discussed at the start, our forecasting framework makes a number of approximations such as neglecting wide-angle effects and correlations between redshift bins. These simplifications (which were instrumental in allowing a direct comparison with galaxy redshift surveys) are only likely to have any material impact on our forecasts at the very largest scales, well away from where the strongest distance constraints come from. As such, the constraints on cosmological parameters that we have presented should not be expected to change appreciably

under a more sophisticated treatment.

More important are the effects that we have accounted for that are sometimes neglected in other forecasts. By explicitly including non-linearities, unknown bias evolution, and foreground subtraction residuals, and systematically exploring parameter degeneracies, we have tried to be as comprehensive (and pessimistic) as possible in acknowledging possible adversities for IM surveys. This should be kept in mind when comparing results from this paper with those from elsewhere. Even so, there is always scope to disagree with the particular decisions that go into any set of forecasts, so we have made our full forecasting code publicly available, with documentation.⁷ The interested reader is encouraged to use it to make their own forecasts.

Acknowledgments. — We are grateful to Y. Akrami, D. Alonso, D. Bacon, T. Baker, B. Bassett, R. Battye, E. Calabrese, T.-C. Chang, H. K. Eriksen, K. Grainge, M. Jones, Y.-C. Li, T. Louis, E. Macaulay, L. Miller, S. Naess, U.-L. Pen, W. Percival, R. Shaw, J. Sievers, K. Sigurdson, O. Smirnov, K. Smith, A. Stebbins, A. Taylor, S. Torchinsky, P. Wilkinson, K. Zarb Adami and J. Zuntz for useful discussions. PB is supported by European Research Council grant StG2010-257080, and acknowledges Oxford Astrophysics, Caltech/JPL, and CITA for hospitality. PGF acknowledges support from Leverhulme, STFC, BIPAC and the Oxford Martin School and the hospitality of the Higgs Centre in Edinburgh. PP is funded by a NRF SKA Postdoctoral Fellowship. MGS acknowledges support from the National Research Foundation (NRF, South Africa), the South African Square Kilometre Array Project and FCT under grant PTDC/FIS-AST/2194/2012.

APPENDIX

HI LINE INTENSITY AND BRIGHTNESS TEMPERATURE

Consider a clump of neutral hydrogen with number density $n_{\text{HI}} = n_0 + n_1$, where 0 and 1 denote the lower and upper level of the hyperfine splitting respectively. The spin temperature, T_S , can be defined using

$$n_1 = n_0 \frac{g_1}{g_0} e^{-\frac{T_*}{T_S}} \simeq 3n_0 = \frac{3}{4}n_{\text{HI}},$$

where $T_* = h\nu_{21}/k_B = 0.0682K$, $g_1 = 3$, $g_0 = 1$, and we have assumed $T_S \gg T_*$. The emissivity (energy per unit time, solid angle/volume, and frequency) of the clump is

$$j_{21} = \frac{A_{10}h\nu_{21}}{4\pi} n_1 \phi(\nu),$$

where $A_{10} \simeq 2.869 \times 10^{-15} \text{ s}^{-1}$ (Wilson et al. 2009) is the Einstein coefficient for spontaneous emission and $\phi(\nu)$ is the line profile, which is assumed to be very narrow, with width $d\nu$ (a simple approximation is $\phi \simeq 1/d\nu$). The clump's luminosity is then

$$dL = \frac{3}{4} A_{10} h\nu_{21} n_{\text{HI}} \phi(\nu) d\nu dA dr,$$

where ν is evaluated in the rest frame of the clump and $dA dr$ is the volume of the clump (dr being the distance along the line of sight) in comoving units, if n_{HI} is the comoving number density. Absorption can be neglected if the spin temperature of the gas is much larger than the background temperature (usually the CMB), so that the total intensity against background just follows from the above luminosity.

The total flux (against the background radiation) from an object at redshift z is then

$$dF = \frac{3h\nu_{21}A_{10}}{16\pi(1+z)^2r^2(z)} n_{\text{HI}} \phi(\nu) d\nu dA dr,$$

⁷<https://gitlab.com/radio-fisher/bao21cm>

where $\nu = \nu_{21}/(1+z)$. Defining the brightness, I , of the clump through $dF \equiv Id\Omega d\nu_{\text{obs}}$ and multiplying by the Rayleigh-Jeans approximation factor, we obtain

$$T_b = \frac{3hc^3 A_{10}}{32\pi k_B \nu_{21}^2} \frac{(1+z)^2}{H(z)} n_{\text{HI}},$$

where we assume that the line width $d\nu/(1+z)$ is much smaller than the observed frequency interval $d\nu_{\text{obs}}$ and the corresponding $dA = r^2 d\Omega$, $dr = \lambda_{21}(1+z)^2/H(z) d\nu_{\text{obs}}$. The comoving number density is

$$n_{\text{HI}} = \Omega_{\text{HI}} \frac{\rho_{c,0}}{m_p} (1 + \delta_{\text{HI}}),$$

where Ω_{HI} is the comoving HI fraction, m_p is the proton mass, δ_{HI} is the HI density contrast, and $\rho_{c,0} = 3H_0^2/8\pi G$ is the critical density today.

REDSHIFT EVOLUTION OF THE HI SIGNAL

In this appendix, we derive the redshift evolution of the HI density, brightness temperature, and bias. We begin by assuming that the HI luminosity from a given spatial volume element (with solid angle $\Delta\Omega$ and frequency interval $\Delta\nu$) is proportional to the HI mass within the volume, M_{HI} . If all the HI contributes to the signal, and the spin temperature is well above the background temperature, the observed brightness temperature of the volume element is

$$T_b(\nu) = \frac{3.23 \times 10^{-4}}{\Delta\Omega \Delta\nu} \frac{M_{\text{HI}}}{(1+z)^2 D_A^2} \text{ mK Mpc}^2 \text{ Hz } M_\odot^{-1},$$

and its proper volume is

$$V_{\text{pix}} = \Delta\Omega \Delta\nu \frac{(c/\nu) D_A^2}{H + dv/ds},$$

where we are taking into account the effect of the peculiar velocity through dv/ds , the proper gradient of the peculiar velocity along the line of sight.

After reionisation, neutral hydrogen can only be found inside galaxies that are able to shield it from ionising UV radiation. The gas temperature and corresponding spin temperature should be much hotter than the CMB, so the approximation above can be used (there are a few cases where a strong background radiation source is capable of generating an absorption signal, but those are negligible given the low resolution we are considering here). We have also neglected HI self-absorption, or any other type of shielding of the HI emission. (Note that even if some re-absorption of the HI signal did happen, we would still expect a linear relation between the HI luminosity and mass, albeit with a smaller constant of proportionality.)

The next step is to connect the HI mass to the underlying halo mass, in order to relate the signal to the cosmological matter density field. We assume that a dark matter halo of mass M contains one or more galaxies with a total mass M_{HI} that is only a function of the halo mass and redshift, i.e. $M_{\text{HI}}(M, z)$. There may be some environmental dependence, which would make this a function of position as well. Some level of stochasticity can also exist in the relation between halo and HI mass, but given the low resolution pixels used in HI intensity mapping experiments, we expect a large number of HI galaxies per pixel, which should average-down any fluctuations and allow us to take the above deterministic relation for the mass function.

To detect the BAO scales at $z = 1$, for example, one needs angular/frequency resolutions of around 1 degree and 5 MHz respectively, which translate into a comoving volume of $1.22 \times 10^5 \text{ Mpc}^3$. In each volume element, we expect a total of around 10^6 dark matter halos with mass between $10^8 - 10^{15} M_\odot$, and $\sim 31,000$ with masses between 5×10^9 and $1 \times 10^{12} M_\odot$ (where the latter range corresponds to halos expected to contain most of the HI mass). This supports our assumption of a position-independent HI mass function due to the averaging over many halos. Some level of stochasticity could still increase the shot noise of the signal, but this is expected to be quite small, as discussed below.

Given $M_{\text{HI}}(M, z)$, we can then relate the signal to the underlying dark matter field. The number of halos of mass M in the observed volume element is given by $[1 + b(M, z)\delta_M(z)] \frac{dn}{dM} dM V_{\text{pix}}$, where δ_M is the underlying dark matter fluctuation at that point in space (and time), $b(M, z)$ is the halo bias, and dn/dM is the proper halo mass function. Integrating over all possible masses, the total observed temperature is then

$$T_b(\nu) = \frac{\alpha}{(1+z)} \frac{\rho_{\text{HI}}(z) [1 + b_{\text{HI}}\delta_M(z)]}{(H + dv/ds)(1 - v/c)},$$

where $\alpha = 2.21 \times 10^{-27} \text{ mK Mpc}^3 M_\odot^{-1} \text{ s}^{-1}$, and the proper HI density and HI bias are

$$\begin{aligned} \rho_{\text{HI}}(z) &= \int_{M_{\text{min}}}^{M_{\text{max}}} dM \frac{dn}{dM} M_{\text{HI}}(M) \\ b_{\text{HI}}(z) &= \rho_{\text{HI}}^{-1} \int_{M_{\text{min}}}^{M_{\text{max}}} dM \frac{dn}{dM} M_{\text{HI}} b(z, M). \end{aligned}$$

Rewriting in terms of the fractional density⁸,

$$\Omega_{\text{HI}}(z) \equiv (1+z)^{-3} \rho_{\text{HI}}(z) / \rho_{c,0},$$

and assuming that the peculiar velocity gradient and v/c terms are small for these large pixels, we finally get

$$T_b(\nu, \Delta\Omega, \Delta\nu) \approx \bar{T}_b(z) \left[1 + b_{\text{HI}} \delta_m(z) - \frac{1}{H(z)} \frac{dv}{ds} \right]$$

$$\bar{T}_b(z) \approx 566h \left(\frac{H_0}{H(z)} \right) \left(\frac{\Omega_{\text{HI}}(z)}{0.003} \right) (1+z)^2 \mu\text{K}.$$

Note that once $M_{\text{HI}}(M, z)$ has been specified, we can calculate Ω_{HI} , the HI bias, and HI brightness temperature in a consistent manner. For the mass function, the most straightforward ansatz is to assume that it is proportional to the halo mass – the constant of proportionality can then be fitted to the available data. Even in this case, however, we need to take into account the fact that not all halos contain galaxies with HI mass. Following Bagla et al. (2010), one can assume that only halos with circular velocities between 30 – 200 km/s are able to host HI, which translates into a halo mass through

$$v_{\text{circ}} = 30\sqrt{1+z} \left(\frac{M}{10^{10} M_{\odot}} \right)^{1/3} \text{ km/s.} \quad (\text{B1})$$

While reasonable, this model is unable to fit constraints on the HI density at high redshift. Possible refinements include allowing the minimum and maximum circular velocities to evolve with redshift, which could make a difference particularly at very low redshifts, or connecting the star formation rate to the halo mass, and then relating that to the HI mass. Alternatively, in Gong et al. (2011a) the relation between HI and halo mass was found using a non-linear function fitted to simulations from Obreschkow et al. (2009) (see their Table 1).

In this paper, we have taken a different approach, using a redshift-independent power-law form for the mass relation,

$$M_{\text{HI}}(M) \propto M^{\alpha}. \quad (\text{B2})$$

An exponent of $\alpha \simeq 0.6$ provides a good fit to current low- and high-redshift constraints when we normalise the relation to the $z = 0.8$ constraints from Switzer et al. (2013). The resulting $\Omega_{\text{HI}}(z)$ is shown in Fig. 20.

Lastly, the shot noise power spectrum due to Poisson fluctuations in halo number is given by

$$P_{\text{HI}}^{\text{shot}}(z) = \left(\frac{\bar{T}_b(z)}{\rho_{\text{HI}}(z)} \right)^2 \int_{M_{\text{min}}}^{M_{\text{max}}} dM \frac{dn}{dM} M_{\text{HI}}^2(M).$$

For the scales we are interested in, this quantity is rather small, although it would be increased somewhat if we allowed some level of stochasticity between each halo and the corresponding HI mass, as described before.

INTERFEROMETER BASELINE DENSITY

In this appendix, we describe how to calculate the baseline density, $n(u)$, for a given array configuration. First of all, one must map-out the uv coverage of the array. For a baseline with position components (L_X, L_Y, L_Z) , the ellipse traced in the uv plane is given by

$$u^2 + \left(v - \frac{L_Z / \lambda \cos \delta_0}{\sin \delta_0} \right)^2 = \frac{L_X^2 + L_Y^2}{\lambda^2},$$

where δ_0 is the declination of the phase tracking centre. For an array with N_d dishes, the total number of unique baselines is $N = N_d(N_d - 1)/2$. Each baseline contributes one elliptical locus, given by the above expression. A full ellipse is traced for each baseline over the course of 24 hours of observation.

Experiment	ν_{min}	ν_{max}	$\delta\nu$ [kHz]	N_d	D_{dish} [m]
ASKAP	700	1800	20	36	12.0
CHIME	400	800	1000	5×256	80×20
JVLA (D)	1000	2000	2000	27	25.0
KAT7	1200	1950	50	7	12.0
MeerKAT	1200	1950	50	64	13.5
SKA1-MID Base	580	1015	50	190	15.0
SKA1-MID Full	580	1015	50	254	14.62
Tianlai	550	950	1000	8×256	100×15

TABLE 5

DETAILS OF THE ARRAY CONFIGURATIONS FOR WHICH $n(u)$ WAS CALCULATED. $\nu_{\text{min,max}}$ ARE IN MHz, AND $\delta\nu$ IS THE CHANNEL BANDWIDTH. THE EFFICIENCY FACTOR FOR ALL DISH ARRAYS WAS TAKEN TO BE $\eta = 0.7$.

⁸The $(1+z)^{-3}$ term shows up here because dn/dm is the halo mass function in proper volume units.




FIG. 30.— Noise sensitivity as a function of transverse wavenumber at $z = 1$, for MeerKAT and SKA1-MID Band 1 (lower is better). The BAO scales (upper plot and grey band) are shown for reference.

The baseline density $n(u)$ is just a histogram of uv coverage in rings centred at the origin, i.e. the number of baselines per ring of radius $|u| = \sqrt{u^2 + v^2}$ and width Δu . The bin size is defined by the field of view (FOV), which depends on the effective area of a single array element, A_e , and the wavelength, λ , of the observation,

$$\Delta u \sim \frac{1}{\sqrt{\text{FOV}}} = \sqrt{A_e/\lambda^2}.$$

We neglect points with $|u| \leq 1/\sqrt{\text{FOV}}$, as these baselines are not independent (see next appendix).

We have computed $n(u)$ for ASKAP, CHIME, JVLA, KAT7, MeerKAT, SKA1-MID, and Tianlai (only some of which have been used in interferometric mode in this paper). For the dish arrays, we generated uv coverages for 24 hour observations, with 60 second integration time per visibility. The uv coverage was scaled depending on the observation frequency, and $n(u)$ computed as described above. As the different arrays operate in different bands, we simulated 10 frequency channels for each, uniformly spaced in each band. Details of the simulated configurations are given in Table 5, and we have made the resulting $n(u)$ available online.⁹ The sensitivity for two of these experiments is shown as a function of transverse wavenumber (at $z = 1$) in Fig. 30. Note that, for the cylinder interferometers CHIME and Tianlai, the effective cylinder length is taken to be smaller than the geometric length, as the cylinders will likely be underilluminated to mitigate edge effects.

DERIVATION OF NOISE EXPRESSIONS

Single-dish (autocorrelation)

For a single dish with effective collecting area A_e , the noise associated with the measured flux is assumed Gaussian with an RMS (flux sensitivity) given by

$$\sigma_S = \frac{2k_B T_{\text{sys}}}{A_e \sqrt{\delta\nu t_p}},$$

integrated over a frequency interval $\delta\nu$ and observation time per pointing t_p . Telescope sensitivities are often quoted in terms of the System Equivalent Flux Density¹⁰, $\text{SEFD} \equiv 2k_B T_{\text{sys}}/A_e$, or alternatively just A_e/T_{sys} . The effective area A_e is usually $\sim 70 - 80\%$ of the actual dish area, depending of the efficiency, η , of the system. The total system temperature is $T_{\text{sys}} = T_{\text{sky}} + T_{\text{inst}}$, where $T_{\text{sky}} \approx 60 (300 \text{ MHz}/\nu)^{2.55}$ K is the sky temperature, and T_{inst} is the instrument temperature (which is typically higher than the sky temperature above 300 MHz). For typical instrumental

⁹<https://gitlab.com/radio-fisher/bao21cm>

¹⁰For a system with many dishes, both SEFD and A_e/T_{sys} are defined in terms of the total collecting area.

specifications, the single-dish noise RMS can be written as

$$\sigma_S = 18 \text{ mJy} \left(\frac{T_{\text{sys}}}{25 \text{ K}} \right) \left(\frac{200 \text{ m}^2}{A_e} \right) \left(\frac{0.1 \text{ MHz}}{\delta\nu} \right)^{1/2} \left(\frac{1 \text{ h}}{t_p} \right)^{1/2}.$$

Because we are interested in brightness temperature sensitivity (i.e. where the signal fills the primary beam), we need to look instead at the intensity of the signal. This is found by dividing the flux by the primary beam solid angle at FWHM, θ_B^2 . In the Rayleigh-Jeans limit, the conversion from intensity to temperature then gives $\sigma_T \approx \lambda^2 T_{\text{sys}} / (\theta_B^2 A_e \sqrt{\delta\nu t_p})$. For a given survey area, $\sim S_{\text{area}} / \theta_B^2$ pointings are needed, so the single-dish RMS noise temperature is

$$\sigma_T \approx \frac{T_{\text{sys}}}{\sqrt{\delta\nu t_{\text{tot}}}} \frac{\lambda^2}{\theta_B^2 A_e} \sqrt{S_{\text{area}} / \theta_B^2},$$

where t_{tot} is the *total* observation time for the survey. For an array with N_d identical dishes, the signals from each dish can be added incoherently, reducing σ_T by a factor of $1/\sqrt{N_d}$. Telescopes can also use focal plane arrays on each dish to increase the instantaneous FOV, effectively increasing the number of beams, N_b , using multiple feeds or Phased Array Feeds (PAFs). Receivers can also support more than one polarisation channel, $n_{\text{pol}} \geq 1$, and the channels can be added incoherently. Taking all of this into account, we can write

$$\sigma_T \approx \frac{T_{\text{sys}}}{\sqrt{n_{\text{pol}} \delta\nu t_{\text{tot}}}} \frac{\lambda^2}{\theta_B^2 A_e} \sqrt{S_{\text{area}} / \theta_B^2} \sqrt{\frac{1}{N_d N_b}},$$

with the constraint that $S_{\text{area}} \geq N_b \theta_B^2$, since nothing is gained by pointing all the feeds in the same direction.

We are interested in a statistical detection of the HI signal. The 3D noise power spectrum associated with an autocorrelation measurement is just $P_N = \sigma_T^2 V_{\text{pix}}$, where $V_{\text{pix}} = (r\theta_B)^2 \times (r_\nu \delta\nu / \nu_{21})$ is the 3D volume of each volume element. We can then obtain the expression for the noise covariance from Eq. (5) (ignoring the beams),

$$C^N(\mathbf{q}, y) \equiv \frac{P_N}{r^2 r_\nu} = \frac{T_{\text{sys}}^2 U_{\text{bin}}}{n_{\text{pol}} t_{\text{tot}} \Delta\nu} \left(\frac{\lambda^4}{A_e^2 \theta_B^4} \right) \mathcal{I},$$

where $U_{\text{bin}} = S_{\text{area}} \Delta\tilde{\nu}$ and $\mathcal{I} = 1/N_b N_d$. For a dish reflector, $A_e \equiv \eta\pi(D_{\text{dish}}/2)^2$ and $\theta_B \approx \lambda/D_{\text{dish}}$, and so the factor in brackets in the expression above is $\mathcal{O}(1)/\eta^2$. For a dish equipped with a PAF, the beams begin to overlap below a critical frequency, ν_{crit} , and so there is a resulting loss of sensitivity, such that

$$C^N(\mathbf{q}, y) \rightarrow C^N(\mathbf{q}, y) \times \begin{cases} 1, & \nu > \nu_{\text{crit}} \\ (\nu_{\text{crit}}/\nu)^2, & \nu \leq \nu_{\text{crit}}. \end{cases} \quad (\text{D1})$$

Interferometer (cross-correlation)

For an interferometer, a pair of elements separated a baseline of length d measures a visibility $V(\mathbf{u}, \nu)$, where \mathbf{u} is the vector in uv space corresponding to that baseline¹¹, and $u = |\mathbf{u}| = d/\lambda$. The uv -space resolution is set by the interferometer FOV, which for an array of dishes is given by the beam solid angle of a single dish, $\delta u \delta v \sim 1/\text{FOV} \sim A_e/\lambda^2$. Visibilities separated by more than this distance in uv -space can be taken as independent.

The detector noise for a single complex visibility measurement, $N(\mathbf{u}, \nu)$, is assumed Gaussian with variance

$$\sigma_T^2 \equiv \langle N(\mathbf{u}, \nu_1) N^*(\mathbf{u}, \nu_2) \rangle = \left(\frac{\lambda^2 T_{\text{sys}}}{A_e \sqrt{\delta\nu t_{\mathbf{u}}}} \right)^2 \delta_{1,2},$$

where $\delta\nu$ is the channel bandwidth and $t_{\mathbf{u}}$ is the observing time for a given baseline. While each visibility measurement is independent in terms of instrumental noise, for the same sky signal the measurements will be strongly correlated for distances smaller than $\sqrt{A_e/\lambda^2}$, as explained above. One way of dealing with this is to average all visibilities falling into each uv -space resolution element of area $\delta u \delta v$. The noise will then be reduced by the number of points in that element, while the sky visibility stays essentially the same (assuming sufficiently high uv resolution).

Let t_s be the integration time for one visibility, and $N_s(\mathbf{u})$ the corresponding total number of visibilities falling into a given element after one “snapshot”. N_s will then be directly related to the baseline distribution (once we have factored in the observation angle), as explained in the previous appendix. The noise in each uv resolution element is then

$$\sigma_T^2(\mathbf{u}, \nu) = \left(\frac{\lambda^2 T_{\text{sys}}}{A_e \sqrt{\delta\nu t_s N_s(\mathbf{u})}} \right)^2.$$

Note that t_s is usually just a few seconds, as longer integration times generate a “smearing” of the visibility in the uv plane due to Earth’s rotation. Smaller integration times allow more visibility points to be sampled, but each with larger noise.

¹¹The projection of that baseline on the plane perpendicular to the line of sight (the telescope pointing) is what actually matters.

The total observation time in a given patch of the sky, t_p , is usually more than the snapshot time t_s , and the telescope tracks the patch. This implies that the same baseline might produce different vectors in the uv plane as the observation angle changes. To allow for different choices of the pixel area and integration time per visibility, one usually refers to the average number density of baselines averaged over a 24h period,

$$n(\mathbf{u}) = N(\mathbf{u})/(\delta u \delta v)/(24h/t_s).$$

$N(\mathbf{u})$ now corresponds to the total number of baselines falling into a given resolution element $\delta u \delta v$ in a 24 hour period, so that sky rotation is taken into account. For a given (square) resolution element, $(\Delta u)^2$, we can write

$$\sigma_T^2(\mathbf{u}, \nu) = \left(\frac{\lambda^2 T_{\text{sys}}}{A_e \sqrt{\delta \nu n(\mathbf{u}) (\Delta u)^2 t_p}} \right)^2,$$

where $n(\mathbf{u})$ is usually only a function of $u \equiv |\mathbf{u}|$ due to the symmetrising effect of the rotation on the uv coverage.

If we assume that $n(u)$ is constant on the uv plane between some $u_{\min} = D_{\min}/\lambda$ and $u_{\max} = D_{\max}/\lambda$ (which is not the same as assuming an uniform distribution of antennas), we can write (see Eq. 6)

$$(\pi u_{\max}^2 - \pi u_{\min}^2) n(u) = N_d(N_d - 1)/2.$$

This follows from noting that the integration of $n(u)$ over the uv plane should give the total number of baselines.

The variance of the noise in 3D Fourier space is related to the variance of the visibilities through

$$\langle N_S(\mathbf{k}) N_S^*(\mathbf{k}) \rangle \approx (\delta \nu / \nu_{21}) \Delta \tilde{\nu} [r^2 r_\nu \sigma_T(\mathbf{u}, \nu)]^2,$$

which is obtained by noting that the 3D Fourier component corresponds to a Fourier transform of the visibility along the frequency direction. The noise power spectrum for a single-pointing observation is then given by

$$P_N = \frac{\langle N_S(\mathbf{k}) N_S^*(\mathbf{k}) \rangle}{(\text{FOV } r^2)(r_\nu \Delta \tilde{\nu})},$$

which, using $(\Delta u)^2 \sim 1/\text{FOV}$, reduces to an expression similar to Eq. (5),

$$C^N(\mathbf{q}, y) \equiv \frac{P_N}{r^2 r_\nu} = \frac{T_{\text{sys}}^2}{\nu_{21} t_p} \frac{\lambda^4}{A_e^2 n(\mathbf{u})}.$$

In the above, t_p is the observation time for a single pointing of the interferometer. Once a signal-to-noise ratio of unity is achieved on the scales of interest, one can gain by moving to another pointing (i.e. increasing the survey area). The time spent at each pointing is then $t_p = t_{\text{tot}}(\text{FOV}/S_{\text{area}})$, and the number of observed modes is increased by a factor of $S_{\text{area}}/\text{FOV}$. Taking this into account, as well as the possibility of having multiple beams, $N_b \geq 1$, and polarisation channels, $n_{\text{pol}} \geq 1$, we arrive at the expression

$$C^N(\mathbf{q}, y) = \frac{T_{\text{sys}}^2}{\nu_{21} n_{\text{pol}} t_{\text{tot}}} \frac{\lambda^4 S_{\text{area}}}{A_e^2 \cdot \text{FOV}} \frac{1}{N_b n(\mathbf{u})}. \quad (\text{D2})$$

For a standard dish reflector, $\text{FOV} \approx \theta_B^2$. For an interferometer equipped with PAFs, the primary beam scales as

$$\theta_B(\nu) = \theta_B(\nu_{\text{crit}}) \times \begin{cases} (\nu_{\text{crit}}/\nu), & \nu > \nu_{\text{crit}} \\ 1, & \nu \leq \nu_{\text{crit}}. \end{cases}$$

For an aperture array, the primary beam scales as usual with frequency ($\theta_B = \theta_B(\nu_{\text{crit}}) \times (\nu_{\text{crit}}/\nu)$), but now the effective area picks up a correction as the array becomes dense below the critical frequency,

$$A_{\text{eff}}(\nu) = A_{\text{eff}}(\nu_{\text{crit}}) \times \begin{cases} (\nu_{\text{crit}}/\nu)^2, & \nu > \nu_{\text{crit}} \\ 1, & \nu \leq \nu_{\text{crit}}. \end{cases}$$

The noise expression for cylinder interferometers (like CHIME) is more complicated. First of all, the primary beam is anisotropic; in the direction along the cylinder, the beam is $\sim 90^\circ$, while it is limited by the cylinder width, w_{cyl} , in the perpendicular direction, giving $\text{FOV} \approx 90^\circ \times \lambda/w_{\text{cyl}}$ (Newburgh et al. 2014). Secondly, many closely-packed feeds illuminate each cylinder, complicating the relationship between number of receivers and collecting area assumed in Eq. D2. To fit the cylinder noise expression into this form, we write $A_e = \eta l_{\text{cyl}} w_{\text{cyl}}/N_{\text{feed}}$, the effective area *per feed*, where N_{feed} is the number of feeds per cylinder, and l_{cyl} is the cylinder length. There is also a restriction on the survey area; the beam cannot be steered and the telescope drift scans, fixing $S_{\text{area}} \sim 30,000 \text{ deg}^2$ (we choose $25,000 \text{ deg}^2$ as an effective area).

Finally, note that most of the quantities above depend on frequency. In particular, the FOV (and thus the minimum angular resolution) changes with frequency, so usually the maximum possible size for the uv -space resolution element is taken. Moreover, the final equation above is an approximation, and the middle of the frequency interval is taken in some of the expressions; otherwise we would need to consider an integral over the frequency when calculating the noise power spectrum.

DERIVATIVES USED IN THE FISHER MATRIX

Most of the derivatives used in the Fisher matrix can be calculated analytically. It is advantageous to use analytic derivatives because their numerical behaviour can be regulated more easily. They can also offer insight into the behaviour of certain constraints.

The kernel of the Fisher integral consists of products of terms of the form

$$\partial_{p_i} \log C^T = (\partial_{p_i} C^S) / C^T,$$

where we recall that $C^T = C^S + C^N + C^F$, and the equality follows from the fact that only the signal covariance, C^S , is a function of the parameters $\{p_i\}$.

Basic parameters— The derivatives for most of the terms in the signal model, Eq. (4), are relatively straightforward:

$$\begin{aligned} \partial_A C^S &= \frac{f_{\text{bao}}(k)}{1 + A f_{\text{bao}}(k)} C^S \\ \partial_b C^S &= \frac{2}{b + f \mu^2} C^S \\ \partial_f C^S &= \frac{2 \mu^2}{b + f \mu^2} C^S \\ \partial_{\sigma_8} C^S &= (2/\sigma_8) C^S \\ \partial_{n_s} C^S &= \log(k/k_{\text{piv}}) C^S \\ \partial_{\sigma_{\text{NL}}^2} C^S &= -k^2 \mu^2 C^S. \end{aligned}$$

We have used the splitting of $P(k)$ into a smooth power spectrum plus a BAO feature from Eq. 14 to calculate the derivative for the BAO amplitude, A . The derivative for σ_8 can be found by renormalising the power spectrum, $P(k) \rightarrow (\sigma_8/\sigma_8^{\text{fid}})^2 P(k)$, and the derivative for n_s by rewriting $P(k)$ in terms of the primordial power spectrum, $P(k) \propto T^2(k) (k/k_{\text{piv}})^{n_s-1}$, where $k_{\text{piv}} = 0.05 \text{ Mpc}^{-1}$.

Distance measures— The derivatives for the distance scales, $\{\alpha_{\perp}, \alpha_{\parallel}\}$, are more complicated, but remain mostly analytic. For each α , the derivative is

$$\partial_{\alpha} C^S = \left[\frac{n_{\alpha}}{\alpha} + \frac{2f}{b + f \mu^2} \partial_{\alpha} \mu^2 + \partial_k \log P(k) \partial_{\alpha} k \right] C^S$$

where $n_{\perp} = 2$, $n_{\parallel} = 1$. Only $\partial_k \log P(k)$ must be evaluated numerically; the other terms are given by

$$\begin{aligned} \partial_{\alpha_{\perp}} \mu^2 &= -2 \alpha_{\perp}^{-1} \chi^2 / (1 + \chi^2)^2 \\ \partial_{\alpha_{\parallel}} \mu^2 &= 2 \alpha_{\parallel}^{-1} \chi^2 / (1 + \chi^2)^2 \\ \partial_{\alpha_{\perp}} k &= \left(\frac{\alpha_{\perp} q}{r} \right)^2 / (k \alpha_{\perp}) \\ \partial_{\alpha_{\parallel}} k &= \left(\frac{\alpha_{\parallel} y}{r_{\nu}} \right)^2 / (k \alpha_{\parallel}), \end{aligned}$$

where we have used the definitions

$$\begin{aligned} \chi &= \frac{\alpha_{\perp} r_{\nu} q}{\alpha_{\parallel} r y} \\ k^2 &= \left(\frac{\alpha_{\perp} q}{r} \right)^2 + \left(\frac{\alpha_{\parallel} y}{r_{\nu}} \right)^2 \\ \mu^2 &= \frac{y^2}{y^2 + \left(\frac{\alpha_{\perp} r_{\nu}}{\alpha_{\parallel} r} \right)^2 q^2}. \end{aligned}$$

Parameters from distance measures— Where we are interested in constraining cosmological parameters rather than functions of redshift, we project from $\{D_A(z), H(z), f(z)\}$ into the parameters $\{h, \Omega_K, \Omega_{\text{DE}}, w_0, w_a, \gamma\}$ using Eq. (20).

To do this, we first assume the following forms for the functions of redshift, based on simple extensions to Λ CDM:

$$\begin{aligned} H(a) &= H_0 \sqrt{\Omega_M a^{-3} + \Omega_{DE}(a) + \Omega_K a^{-2}} \\ r(a) &= \frac{c}{H_0} S \left(\int \frac{da'}{a'^2 E(a')} \right) \\ f(a) &= \Omega_M^\gamma(a), \end{aligned}$$

where for $\Omega_K = \{-ve, 0, +ve\}$ we have defined $S(x) = \left\{ \sin(x\sqrt{|\Omega_K|})/\sqrt{|\Omega_K|}, x, \sinh(x\sqrt{\Omega_K})/\sqrt{\Omega_K} \right\}$, and

$$\begin{aligned} w(a) &\approx w_0 + (1-a)w_a \\ \Omega_M(a) &= H_0^2 \Omega_M a^{-3} / H^2(a) \\ \Omega_{DE}(a) &= \Omega_{DE} \frac{\exp[3w_a(a-1)]}{a^{3(1+w_0+w_a)}}, \end{aligned}$$

$H(a) \equiv H_0 E(a)$, $H_0 = 100 h \text{ kms}^{-1} \text{ Mpc}^{-1}$, and $\Omega_M = 1 - \Omega_K - \Omega_{DE}$.

Next, we need the derivatives of the original functions of redshift with respect to the new parameters. Most of them enter through the dimensionless Hubble rate, $E(a)$, for which the relevant derivatives are

$$\begin{aligned} \partial_{\Omega_K} E(a) &= (a^{-2} - a^{-3})/2E(a) \\ \partial_{\Omega_{DE}} E(a) &= (1 - a^{-3})/2E(a) \\ \partial_{w_0} E(a) &= -\frac{3}{2} \frac{\Omega_{DE}}{E(a)} \log a \\ \partial_{w_a} E(a) &= -\frac{3}{2} \frac{\Omega_{DE}}{E(a)} [\log a + (1-a)] \\ \partial_\gamma E(a) &= \partial_h E(a) = 0. \end{aligned}$$

For α_\parallel , we then have (with β being any of the parameters except h , and evaluating on Λ CDM)

$$\partial_\beta \alpha_\parallel = \frac{\partial_\beta E}{E_\Lambda}.$$

The expression for α_\perp is more complicated. For all but Ω_K , the derivatives are given by

$$\begin{aligned} \partial_\beta \alpha_\perp &= -\frac{\alpha_\perp}{r_\Lambda} \frac{\partial S(x)}{\partial x} \partial_\beta x \\ \partial_\beta x &= -\int \frac{c da'}{a'^2 E^2(a')} \partial_\beta E(a'). \end{aligned}$$

Additional terms appear in the derivative w.r.t. Ω_K ,

$$\partial_{\Omega_K} \alpha_\perp = -\frac{\alpha_\perp}{r_\Lambda} \left[\frac{\partial S(x)}{\partial x} \partial_{\Omega_K} x + \frac{\Theta(\Omega_K)}{2\Omega_K} \left(x \frac{\partial S(x)}{\partial x} - S(x) \right) \right],$$

where $\Theta = 0$ for $\Omega_K = 0$ and unity elsewhere. For both α_\parallel and α_\perp , the h derivative is $\partial_h \alpha = \alpha/h$.

For the growth function, the derivatives for all but $\{\gamma, \Omega_K, \Omega_{DE}\}$ are given by

$$\partial_\beta f(a) = -\gamma f(a) \frac{\partial_\beta E}{E_\Lambda}.$$

The derivatives with respect to the other parameters are

$$\begin{aligned} \partial_\gamma f &= f(a) \log \Omega_M(a) \\ \partial_{\Omega_K} f &= -\gamma f [\Omega_M^{-1} + \partial_{\Omega_K} \log E] \\ \partial_{\Omega_{DE}} f &= -\gamma f [\Omega_M^{-1} + \partial_{\Omega_{DE}} \log E]. \end{aligned}$$

REFERENCES

- Abramo, L. R., & Leonard, K. E. 2013, Mon.Not.Roy.Astron.Soc., 432, 318, [arXiv:1302.5444]
 Albrecht, A., Amendola, L., Bernstein, G., Clowe, D., Eisenstein, D., et al. 2009, [arXiv:0901.0721]
 Albrecht, A., Bernstein, G., Cahn, R., Freedman, W. L., Hewitt, J., et al. 2006, [arXiv:astro-ph/0609591]
 Alcock, C., & Paczynski, B. 1979, Nature, 281, 358
 Alonso, D., Ferreira, P. G., & Santos, M. G. 2014, [arXiv:1405.1751]
 Amendola, L., Appleby, S., Bacon, D., et al. 2013, Living Reviews in Relativity, 16, 6, [arXiv:1206.1225]

- Ansari, R., Campagne, J. E., Colom, P., et al. 2012, *A. & A.*, 540, A129, [arXiv:1108.1474]
- ATNF. 2014, Parkes Radio Telescope Users Guide, http://www.parkes.atnf.csiro.au/observing/documentation/user_guide/pks_ug.html
- Austermann, J., Aird, K., Beall, J., Becker, D., Bender, A., et al. 2012, *Proc.SPIE Int.Soc.Opt.Eng.*, 8452, 84520E, [arXiv:1210.4970]
- Bagla, J., Khandai, N., & Datta, K. K. 2010, *Mon. Not. Roy. Astron. Soc.*, 407, 567, [arXiv:0908.3796]
- Baker, T., Ferreira, P. G., & Skordis, C. 2013, *Phys.Rev.*, D87, 024015, [arXiv:1209.2117]
- Barkana, R., & Loeb, A. 2005, *Ap. J. L.*, 624, L65, [arXiv:astro-ph/0409572]
- Battye, R. A., Browne, I. W. A., Dickinson, C., Heron, G., Maffei, B., & Pourtsidou, A. 2013, *M.N.R.A.S.*, 434, 1239, [arXiv:1209.0343]
- Battye, R. A., Davies, R. D., & Weller, J. 2004, *Mon. Not. Roy. Astron. Soc.*, 355, 1339, [arXiv:astro-ph/0401340]
- Battye, R. A., & Pearson, J. A. 2012, *JCAP*, 1207, 019, [arXiv:1203.0398]
- Beck, R. 2001, *Sp. Sci. Rev.*, 99, 243, [arXiv:astro-ph/0012402]
- Blake, C., & Glazebrook, K. 2003, *Ap.J.*, 594, 665, [arXiv:astro-ph/0301632]
- Braun, R., Thilker, D., & Walterbos, R. A. M. 2003, *A & A*, 406, 829, [arXiv:astro-ph/0305403]
- Bull, P., & Kamionkowski, M. 2013, *Phys. Rev. D*, 87, 081301, [arXiv:1302.1617]
- Camera, S., Santos, M. G., Ferreira, P. G., & Ferramacho, L. 2013, *Phys.Rev.Lett.*, 111, 171302, [arXiv:1305.6928]
- Chang, T.-C., & GBT-HIM Team. 2014, in *Exascale Radio Astronomy*, 50102
- Chang, T.-C., Pen, U.-L., Bandura, K., & Peterson, J. B. 2010, *Nature*, 466, 463, [arXiv:1007.3709]
- Chang, T.-C., Pen, U.-L., Peterson, J. B., & McDonald, P. 2008, *Phys. Rev. Lett.*, 100, 091303, [arXiv:0709.3672]
- Chen, X. 2012, *International Journal of Modern Physics Conference Series*, 12, 256, [arXiv:1212.6278]
- CHIME Collaboration. 2012, CHIME Overview, http://chime.phas.ubc.ca/CHIME_overview.pdf
- Clarkson, C., Cort  s, M., & Bassett, B. 2007, *JCAP*, 8, 11, [arXiv:astro-ph/0702670]
- Contreras, C., et al. 2013, [arXiv:1302.5178]
- Crocce, M., & Scoccimarro, R. 2008, *Phys.Rev.*, D77, 023533, [arXiv:0704.2783]
- Dewdney, P. E., et al. 2013, SKA Organisation
- Duffy, A. R., Kay, S. T., Battye, R. A., et al. 2012, *M.N.R.A.S.*, 420, 2799, [arXiv:1107.3720]
- Efstathiou, G., & Bond, J. 1999, *Mon.Not.Roy.Astron.Soc.*, 304, 75, [arXiv:astro-ph/9807103]
- Feldman, H. A., Kaiser, N., & Peacock, J. A. 1994, *Ap.J.*, 426, 23, [arXiv:astro-ph/9304022]
- Ferreira, P. G., & Jaffe, A. H. 2000, *M.N.R.A.S.*, 312, 89, [arXiv:astro-ph/9909250]
- Frieman, J., & Dark Energy Survey Collaboration. 2013, in *American Astronomical Society Meeting Abstracts*, Vol. 221, 335.01
- Gleser, L., Nusser, A., & Benson, A. J. 2008, *M.N.R.A.S.*, 391, 383, [arXiv:0712.0497]
- Gong, Y., Chen, X., Silva, M., Cooray, A., & Santos, M. G. 2011a, *ApJ*, 740, L20, [arXiv:1108.0947]
- Gong, Y., Cooray, A., Silva, M. B., Santos, M. G., & Lubin, P. 2011b, *Ap.J.L.*, 728, L46, [arXiv:1101.2892]
- Guth, A. H., & Nomura, Y. 2012, *Phys.Rev.*, D86, 023534, [arXiv:1203.6876]
- Hall, A., Bonvin, C., & Challinor, A. 2013, *Phys.Rev.*, D87, 064026, [arXiv:1212.0728]
- Hawken, A. J., Abdalla, F. B., H  tsi, G., & Lahav, O. 2012, *M.N.R.A.S.*, 424, 2, [arXiv:1111.2544]
- Hogg, D. W. 1999, *ArXiv*, [arXiv:astro-ph/9905116]
- Holdaway, M. A. 1999, in *Astronomical Society of the Pacific Conference Series*, Vol. 180, *Synthesis Imaging in Radio Astronomy II*, ed. G. B. Taylor, C. L. Carilli, & R. A. Perley, 401
- Hu, W., & Sawicki, I. 2007, *Phys. Rev. D*, 76, 104043
- Ishak, M., & Dossett, J. 2009, *Phys.Rev.*, D80, 043004, [arXiv:0905.2470]
- Jeli  , V., et al. 2008, *M.N.R.A.S.*, 389, 1319, [arXiv:0804.1130]
- Johnston, S., Taylor, R., Bailes, M., et al. 2008, *Experimental Astronomy*, 22, 151, [arXiv:0810.5187]
- Jonas, J. L. 2009, *IEEE Proceedings*, 97, 1522
- Kaiser, N. 1987, *Mon.Not.Roy.Astron.Soc.*, 227, 1
- Kerp, J., Winkel, B., Ben Bekhti, N., Fl  er, L., & Kalberla, P. M. W. 2011, *Astronomische Nachrichten*, 332, 637, [arXiv:1104.1185]
- Khandai, N., Sethi, S. K., Di Matteo, T., et al. 2011, *Mon.Not.Roy.Astron.Soc.*, 415, 2580, [arXiv:1012.1880]
- Kleban, M., & Schillo, M. 2012, *JCAP*, 1206, 029, [arXiv:1202.5037]
- Lah, P., Chengalur, J. N., Briggs, F. H., Colless, M., De Propriis, R., et al. 2007, *Mon.Not.Roy.Astron.Soc.*, 376, 1357, [arXiv:astro-ph/0701668]
- Lewis, A., Challinor, A., & Lasenby, A. 2000, *Astrophys. J.*, 538, 473, [arXiv:astro-ph/9911177]
- Li, C., Jing, Y., Kauffmann, G., Boerner, G., Kang, X., et al. 2007, *Mon.Not.Roy.Astron.Soc.*, 376, 984, [arXiv:astro-ph/0701218]
- Lidz, A., Furlanetto, S. R., Oh, S. P., et al. 2011, *ApJ*, 741, 70, [arXiv:1104.4800]
- Linder, E. V. 2005, *Phys.Rev.*, D72, 043529, [arXiv:astro-ph/0507263]
- Linder, E. V., & Cahn, R. N. 2007, *Astropart.Phys.*, 28, 481, [arXiv:astro-ph/0701317]
- Liu, A., Parsons, A. R., & Trott, C. M. 2014, *ArXiv e-prints*, [arXiv:1404.2596]
- Liu, A., & Tegmark, M. 2011, *Phys. Rev. D*, 83, 103006, [arXiv:1103.0281]
- Liu, A., Tegmark, M., Bowman, J., et al. 2009, *M.N.R.A.S.*, 398, 401, [arXiv:0903.4890]
- Loeb, A., & Wyithe, J. S. B. 2008, *Physical Review Letters*, 100, 161301, [arXiv:0801.1677]
- Mangum, J. G., Emerson, D. T., & Greisen, E. W. 2007, *A&A*, 474, 679, [arXiv:0709.0553]
- Mao, Y., Tegmark, M., McQuinn, M., Zaldarriaga, M., & Zahn, O. 2008, *Phys. Rev. D*, 78, 023529, [arXiv:0802.1710]
- Martin, A. M., Papastergis, E., Giovanelli, R., Haynes, M. P., Springob, C. M., et al. 2010, *Astrophys.J.*, 723, 1359, [arXiv:1008.5107]
- Masui, K. W., Schmidt, F., Pen, U.-L., & McDonald, P. 2010, *Phys.Rev.*, D81, 062001, [arXiv:0911.3552]
- Masui, K. W., et al. 2013, *Ap. J. L.*, 763, L20, [arXiv:1208.0331]
- McDonald, P., & Seljak, U. 2009, *JCAP*, 0910, 007, [arXiv:0810.0323]
- McQuinn, M., Zahn, O., Zaldarriaga, M., et al. 2006, *ApJ*, 653, 815, [arXiv:astro-ph/0512263]
- Meiring, J. D., Tripp, T. M., Prochaska, J. X., Tumlinson, J., Werk, J., et al. 2011, *Astrophys.J.*, 732, 35, [arXiv:1102.3927]
- MFAA. 2014, Mid-Frequency Aperture Array, <https://www.skatelescope.org/mfaa/>
- Morales, M. F., Bowman, J. D., & Hewitt, J. N. 2006, *ApJ*, 648, 767, [arXiv:astro-ph/0510027]
- Morales, M. F., Hazelton, B., Sullivan, I., & Beardsley, A. 2012, *ApJ*, 752, 137, [arXiv:1202.3830]
- Mortonson, M. J. 2009, *Phys.Rev.*, D80, 123504, [arXiv:0908.0346]
- Napier, P. J., et al. 1994, *Proceedings of the IEEE*, 82, 658
- Newburgh, L. B. et al. 2014, in *Society of Photo-Optical Instrumentation Engineers (SPIE) Conference Series*, Vol. 9145, *Society of Photo-Optical Instrumentation Engineers (SPIE) Conference Series*, 4, [arXiv:1406.2267]
- Niemack, M., Ade, P., Aguirre, J., Barrientos, F., Beall, J., et al. 2010, *Proc.SPIE Int.Soc.Opt.Eng.*, 7741, 77411S, [arXiv:1006.5049]
- Noterdaeme, P., Petitjean, P., Ledoux, C., & Srianand, R. 2009, *RAO*, 0908.1574
- NRAO. 2014, VLA Observational Status Summary 2014B, <https://science.nrao.edu/facilities/vla/docs/manuals/oss>
- Obreschkow, D., Kl  ckner, H.-R., Heywood, I., Levrier, F., & Rawlings, S. 2009, *ApJ*, 703, 1890, [arXiv:0908.0983]
- Oh, S. P., & Mack, K. J. 2003, *M.N.R.A.S.*, 346, 871, [arXiv:astro-ph/0302099]
- Okouma, P. M., Fantaye, Y., & Bassett, B. A. 2013, *Physics Letters B*, 719, 1, [arXiv:1207.3000]
- Oosterloo, T., Verheijen, M., & van Cappellen, W. 2010, *PoS, ISKAF2010*, 043, [arXiv:1007.5141]
- Padmanabhan, H., Choudhury, T. R., & Refregier, A. 2014, [arXiv:1407.6366]
- Parsons, A. R., Pober, J. C., Aguirre, J. E., Carilli, C. L., Jacobs, D. C., & Moore, D. F. 2012, *ApJ*, 756, 165, [arXiv:1204.4749]
- Peebles, P. 1980, *The large-scale structure of the universe* (Princeton University Press)
- Peterson, J. B., Aleksan, R., Ansari, R., et al. 2009, in *ArXiv Astrophysics e-prints*, [arXiv:0902.3091]
- Petrovic, N., & Oh, S. P. 2011, *M.N.R.A.S.*, 413, 2103, [arXiv:1010.4109]
- Planck Collaboration. 2013a, [arXiv:1303.5072]
- . 2013b, *ArXiv e-prints*, [arXiv:1303.5076]
- Pober, J. C., Parsons, A. R., DeBoer, D. R., et al. 2013, *Ap. J.*, 145, 65, [arXiv:1210.2413]
- Pritchard, J. R., & Loeb, A. 2008, *Phys. Rev. D*, 78, 103511, [arXiv:0802.2102]

- Prochaska, J. X., & Wolfe, A. M. 2009, *ApJ*, 696, 1543, [arXiv:0811.2003]
- Rao, S. M., Turnshek, D. A., & Nestor, D. 2006, *Astrophys.J.*, 636, 610, [arXiv:astro-ph/0509469]
- Rassat, A., Amara, A., Amendola, L., Castander, F. J., Kitching, T., et al. 2008, [arXiv:0810.0003]
- Reid, B. A., Samushia, L., White, M., et al. 2012, [arXiv:1203.6641]
- Reid, B. A., & White, M. 2011, *Mon.Not.Roy.Astron.Soc.*, 417, 1913, [arXiv:1105.4165]
- Samushia, L., Percival, W. J., Guzzo, L., et al. 2011, *M.N.R.A.S.*, 410, 1993, [arXiv:1006.0609]
- Samushia, L., Percival, W. J., & Raccanelli, A. 2012, *Mon.Not.Roy.Astron.Soc.*, 420, 2102, [arXiv:1102.1014]
- Sánchez, E., Alonso, D., Sánchez, F. J., García-Bellido, J., & Sevilla, I. 2013, *M.N.R.A.S.*, 434, 2008, [arXiv:1210.6446]
- Santos, M. G., Cooray, A., & Knox, L. 2005, *ApJ*, 625, 575, [arXiv:astro-ph/0408515]
- Seo, H.-J., Dodelson, S., Marriner, J., McGinnis, D., Stebbins, A., et al. 2010, *Astrophys.J.*, 721, 164, [arXiv:0910.5007]
- Shafieloo, A., & Linder, E. V. 2011, *Phys.Rev.*, D84, 063519, [arXiv:1107.1033]
- Shaw, J. R., Sigurdson, K., Sitwell, M., Stebbins, A., & Pen, U.-L. 2014, *Astrophys.J.*, [arXiv:1401.2095]
- SKA South Africa. 2014, KAT-7, <http://public.ska.ac.za/kat-7>
- Smith, R., et al. 2003, *Mon.Not.Roy.Astron.Soc.*, 341, 1311, [arXiv:astro-ph/0207664]
- Smith, R. E., Scoccimarro, R., & Sheth, R. K. 2008, *Phys.Rev.*, D77, 043525, [arXiv:astro-ph/0703620]
- Smoot, G. F., & Debono, I. 2014, [arXiv:1407.3583]
- Swarup, G., et al. 1991, *Current Science*, 60, 95
- Switzer, E., Masui, K., Bandura, K., Calin, L. M., Chang, T. C., et al. 2013, [arXiv:1304.3712]
- Tegmark, M., Hamilton, A. J. S., Strauss, M. A., Vogeley, M. S., & Szalay, A. S. 1998, *Ap.J.*, 499, 555, [arXiv:astro-ph/9708020]
- Vardanyan, M., Trotta, R., & Silk, J. 2009, *M.N.R.A.S.*, 397, 431, [arXiv:0901.3354]
- Wang, X., Tegmark, M., Santos, M. G., & Knox, L. 2006, *ApJ*, 650, 529, [arXiv:astro-ph/0501081]
- White, M., Carlstrom, J. E., Dragovan, M., & Holzapfel, W. L. 1999, *ApJ*, 514, 12, [arXiv:astro-ph/9712195]
- White, M., Song, Y.-S., & Percival, W. J. 2008, *Mon.Not.Roy.Astron.Soc.*, 397, 1348, [arXiv:0810.1518]
- Wilman, R. J. et al. 2008, *M.N.R.A.S.*, 388, 1335, [arXiv:0805.3413]
- Wilson, T. L., Rohlfs, K., & Hüttemeister, S. 2009, *Tools of radio astronomy*, Vol. 86 (Springer)
- Wolz, L., Abdalla, F. B., Blake, C., Shaw, J. R., Chapman, E., & Rawlings, S. 2013, *ArXiv e-prints*, [arXiv:1310.8144]
- Wolz, L., Kilbinger, M., Weller, J., & Giannantonio, T. 2012, *J.C.A.P.*, 9, 9, [arXiv:1205.3984]
- Wyithe, J. S. B., & Loeb, A. 2008, *M.N.R.A.S.*, 383, 606
- Wyithe, J. S. B., Loeb, A., & Geil, P. M. 2008, *M.N.R.A.S.*, 383, 1195
- Zhang, P. 2008, [arXiv:0802.2416]
- Zwaan, M. A., Meyer, M., Staveley-Smith, L., & Webster, R. 2005, *Mon.Not.Roy.Astron.Soc.*, 359, L30, [arXiv:astro-ph/0502257]